

Introduction to High Performance Reactive Architectures & Spring WebFlux

Disclaimer

All information presented in this document and all supplementary materials and programming code represent only my personal opinion and current understanding and has not received any endorsement or approval by IPT - Intellectual Products and Technologies or any third party. It should not be taken as any kind of advice, and should not be used for making any kind of decisions with potential commercial impact.

The information and code presented may be incorrect or incomplete. It is provided "as is", without warranty of any kind, express or implied, including but not limited to the warranties of merchantability, fitness for a particular purpose and non-infringement. In no event shall the author or copyright holders be liable for any claim, damages or other liability, whether in an action of contract, tort or otherwise, arising from, out of or in connection with the information, materials or code presented or the use or other dealings with this information or programming code.

About Me



Trayan Iliev

- CEO of IPT – [Intellectual Products & Technologies](#)
- Oracle® certified programmer [15+ Y](#)
- End-to-end reactive fullstack apps with [Go](#), [Python](#), [Java](#), [ES6/7](#), [TypeScript](#), [Angular](#), [React](#) and [Vue.js](#)
- [12+ years](#) IT trainer
- [Voxxed Days](#), [jPrime](#), [Java2Days](#), [jProfessionals](#), [BGOUG](#), [BGJUG](#), [DEV.BG](#) speaker
- Lecturer @ [Sofia University](#) – courses: Fullstack React, Multimedia with Angular & TypeScript, Spring 5 Reactive, Internet of Things (with SAP), Multiagent Systems and Social Robotics, Practical Robotics & Smart Things
- [Robotics / smart-things/ IoT](#) enthusiast, [RoboLearn](#) hackathons organizer

Where to Find The Code and Materials?

<https://github.com/iproduct/course-ml>

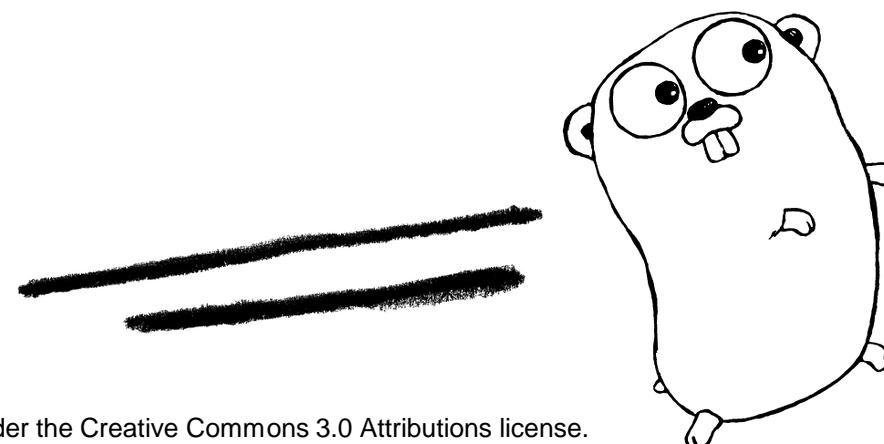
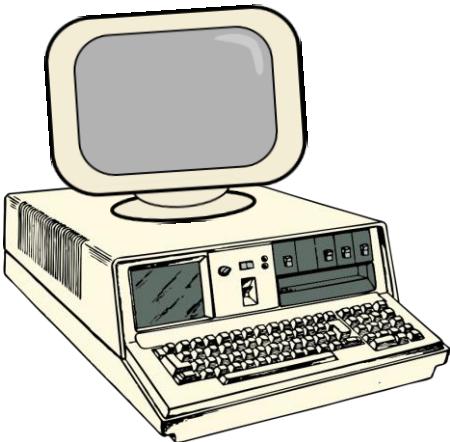
Machine Learning + Big Data in Real Time +
Cloud Technologies

=> The Future of Intelligent Systems

Asynchronous Event Streams in Real-time

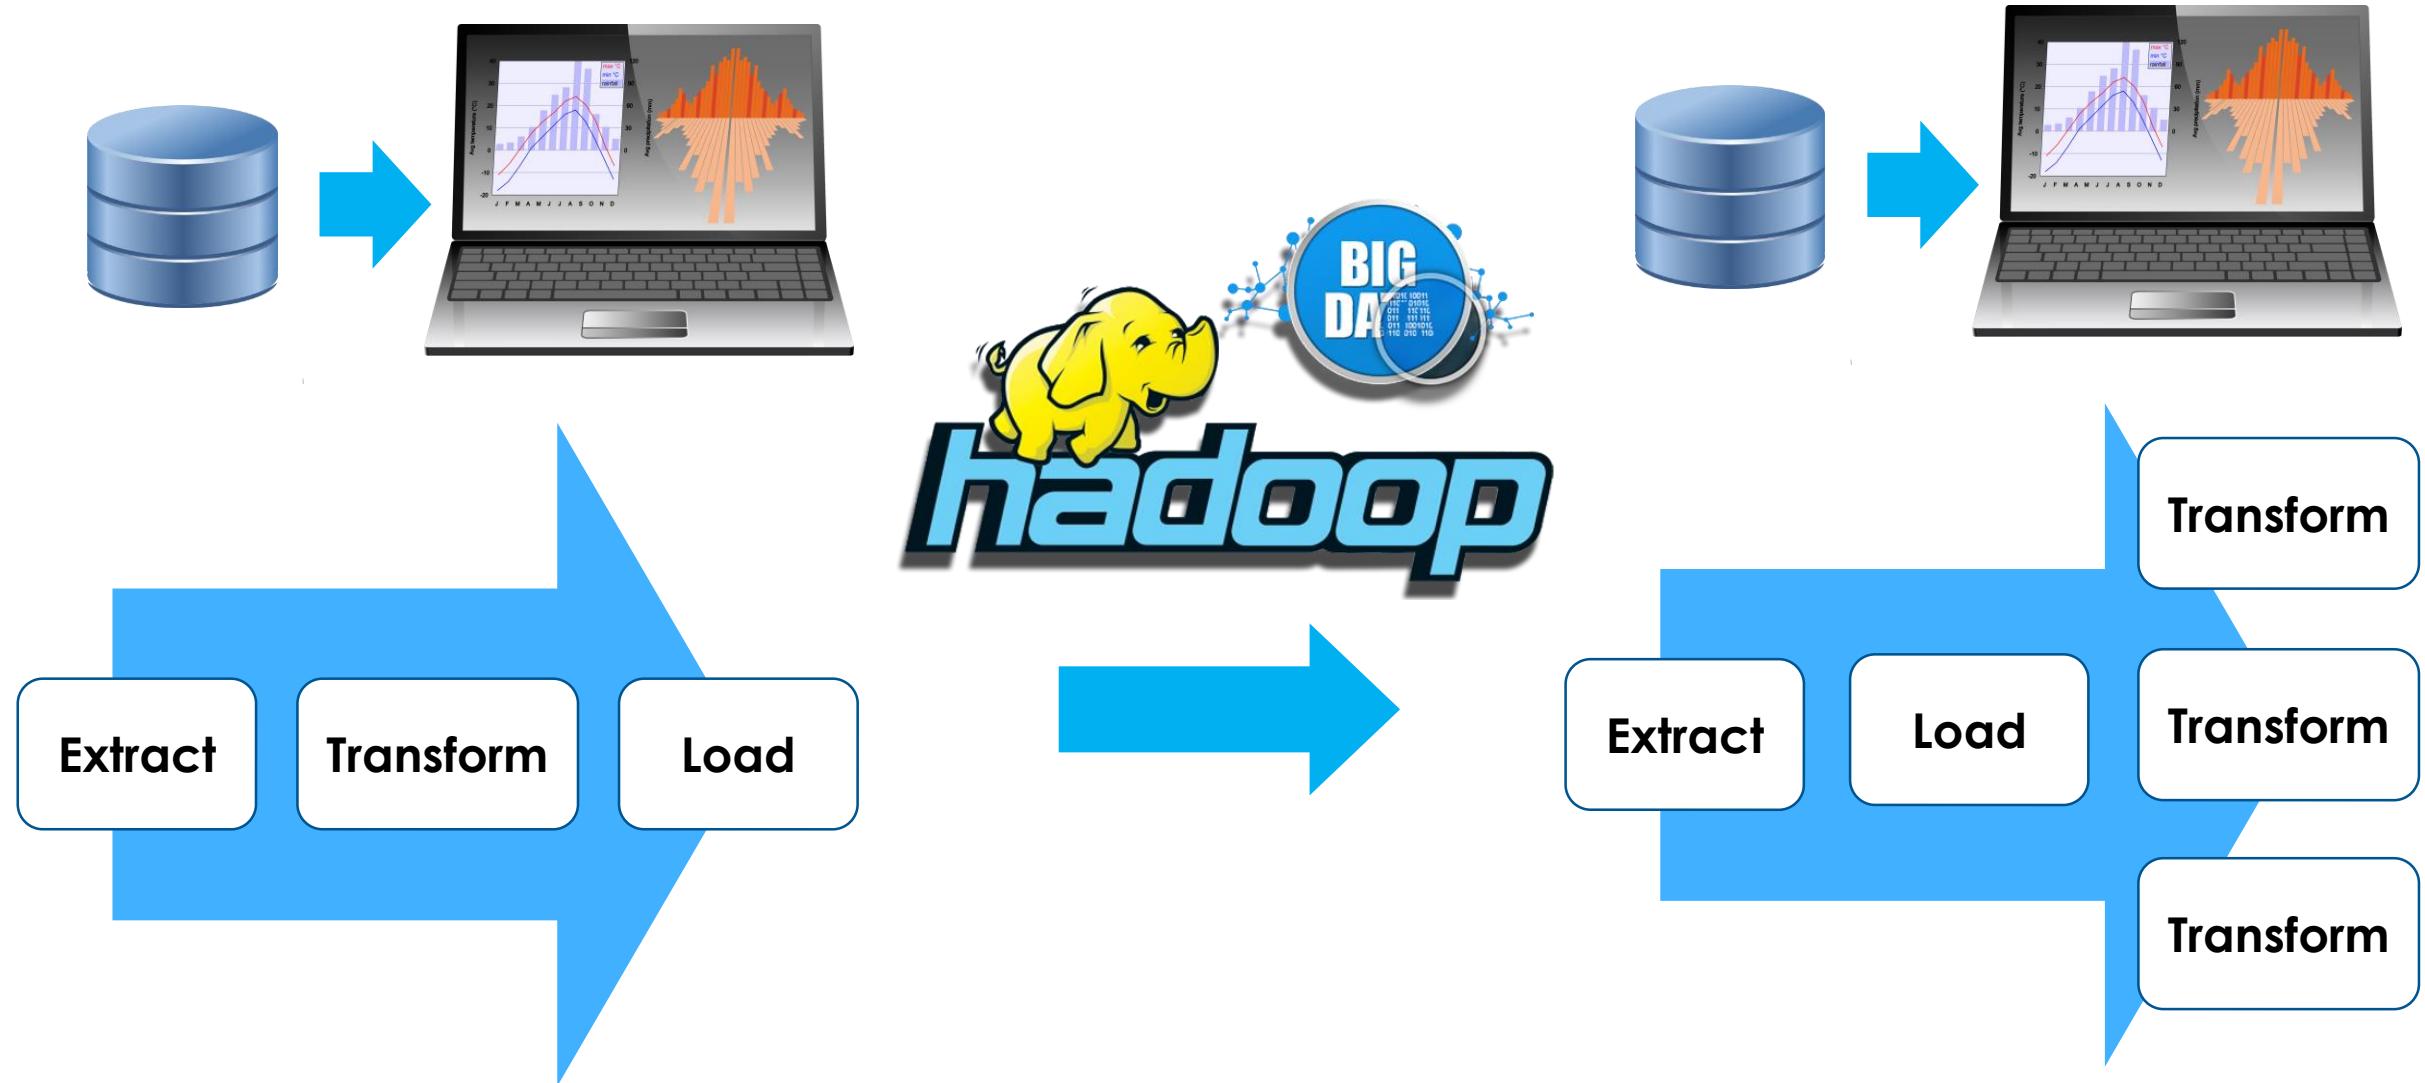


Need for Speed :)



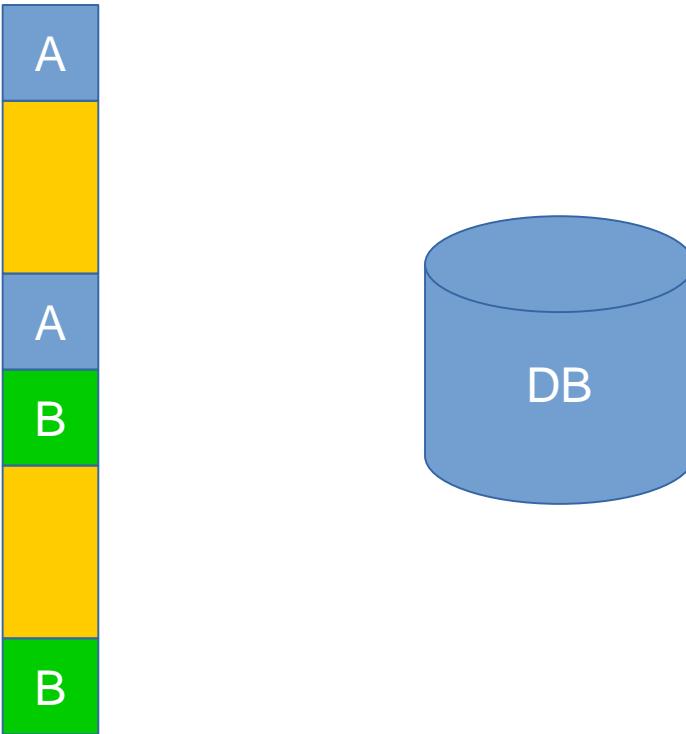
The Go gopher was designed by Renee French. (<http://reneefrench.blogspot.com/>) The design is licensed under the Creative Commons 3.0 Attributions license.

Batch Processing

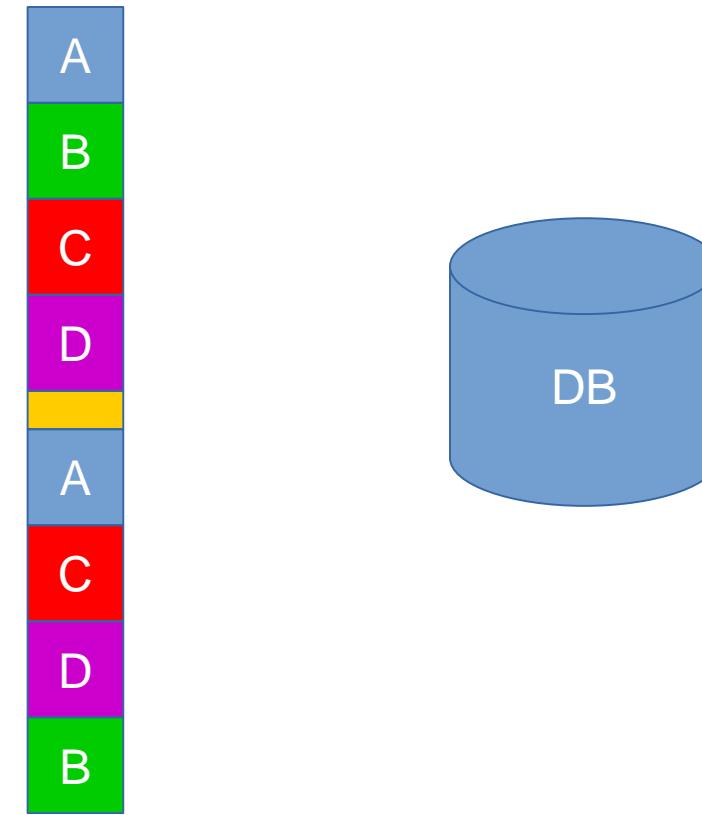


Synchronous vs. Asynchronous IO

Synchronous

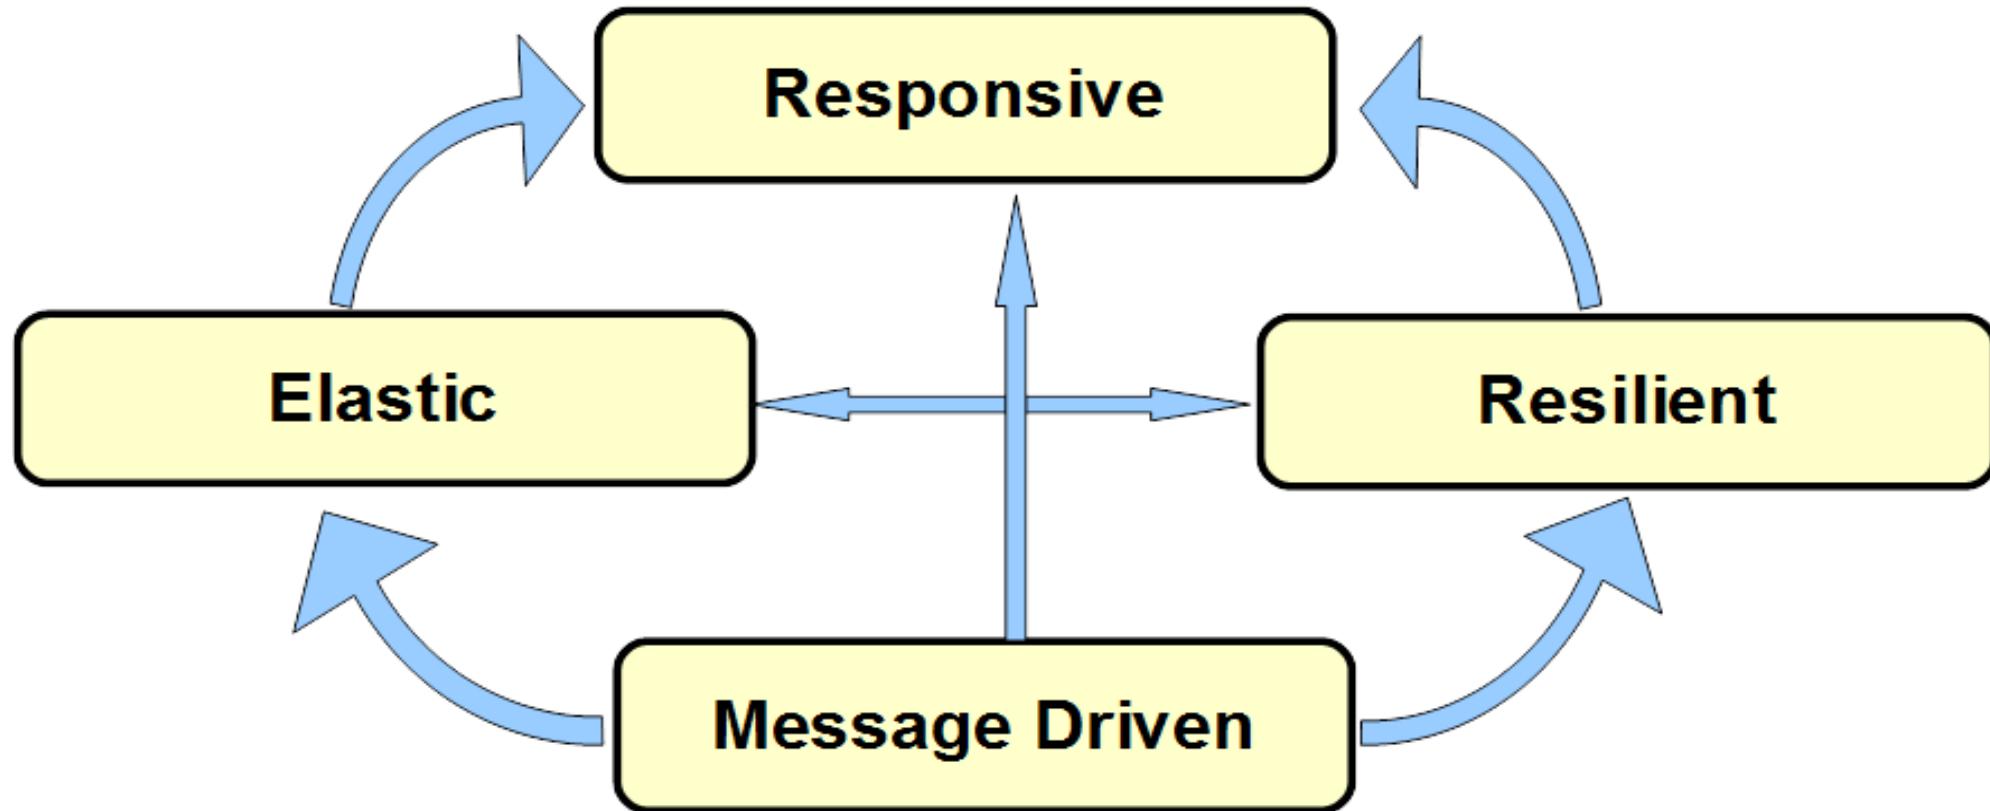


Asynchronous



Reactive Manifesto

<http://www.reactivemanifesto.org>



Scalable, Massively Concurrent

- **Message Driven** – asynchronous message-passing allows to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages [Reactive Manifesto].
- The main idea is to separate concurrent producer and consumer workers by using **message queues**.
- **Message queues** can be **unbounded** or **bounded** (limited max number of messages)
- **Unbounded** message queues can present memory allocation problem in case the producers outrun the consumers for a long period → **OutOfMemoryError**

What's High Performance?

- ❖ **Performance** is about 2 things (Martin Thompson – <http://www.infoq.com/articles/low-latency-vp>):
 - **Throughput** – units per second, and
 - **Latency** – response time
- ❖ **Real-time** – time constraint from input to response regardless of system load.
- ❖ **Hard real-time system** if this constraint is not honored then a total system failure can occur.
- ❖ **Soft real-time system** – low latency response with little deviation in response time
- ❖ **100 nano-seconds to 100 milli-seconds.** [Peter Lawrey]

Data / Event / Message Streams

“Conceptually, a stream is a (potentially never-ending) **flow of data records**, and a transformation is an operation that takes one or more streams as input, and produces one or more output streams as a result.”

Apache Flink: Dataflow Programming Model

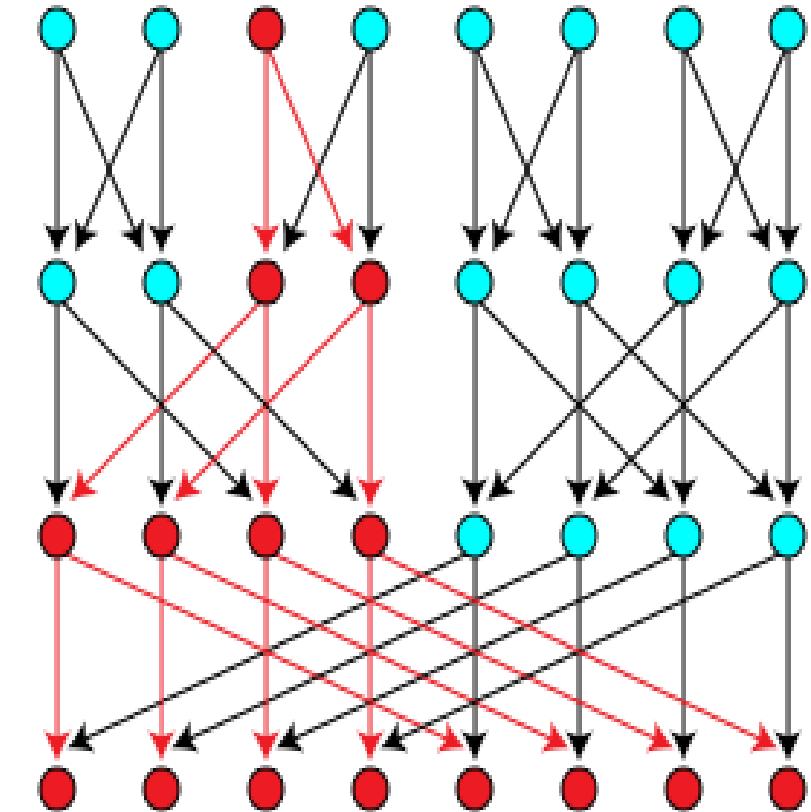
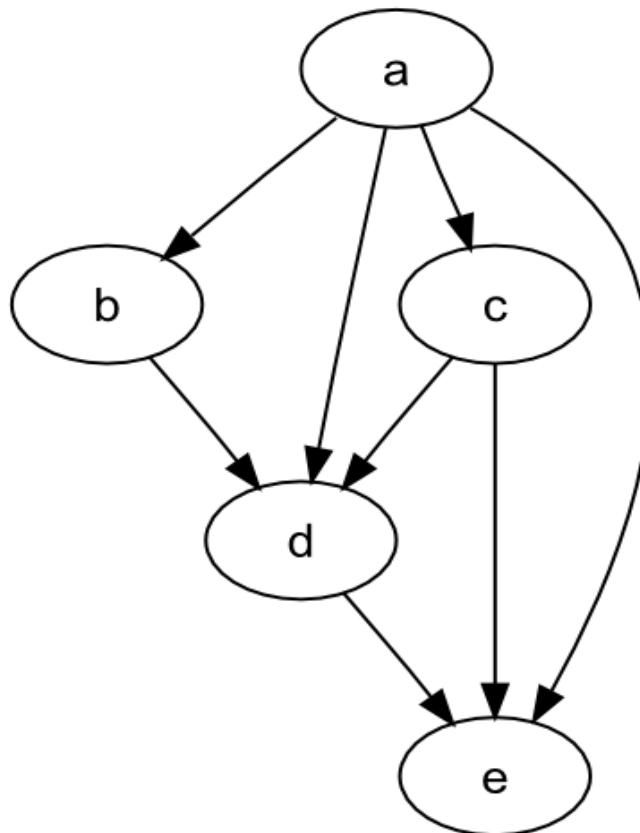
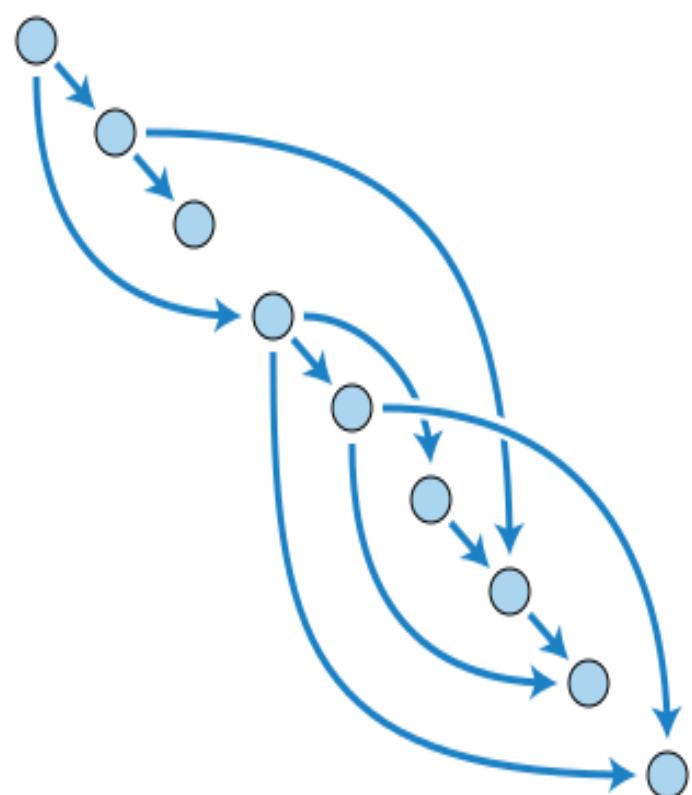
Data Stream Programming

The idea of **abstracting logic from execution** is hardly new -- it was the dream of **SOA**. And the recent emergence of **microservices** and **containers** shows that the dream still lives on.

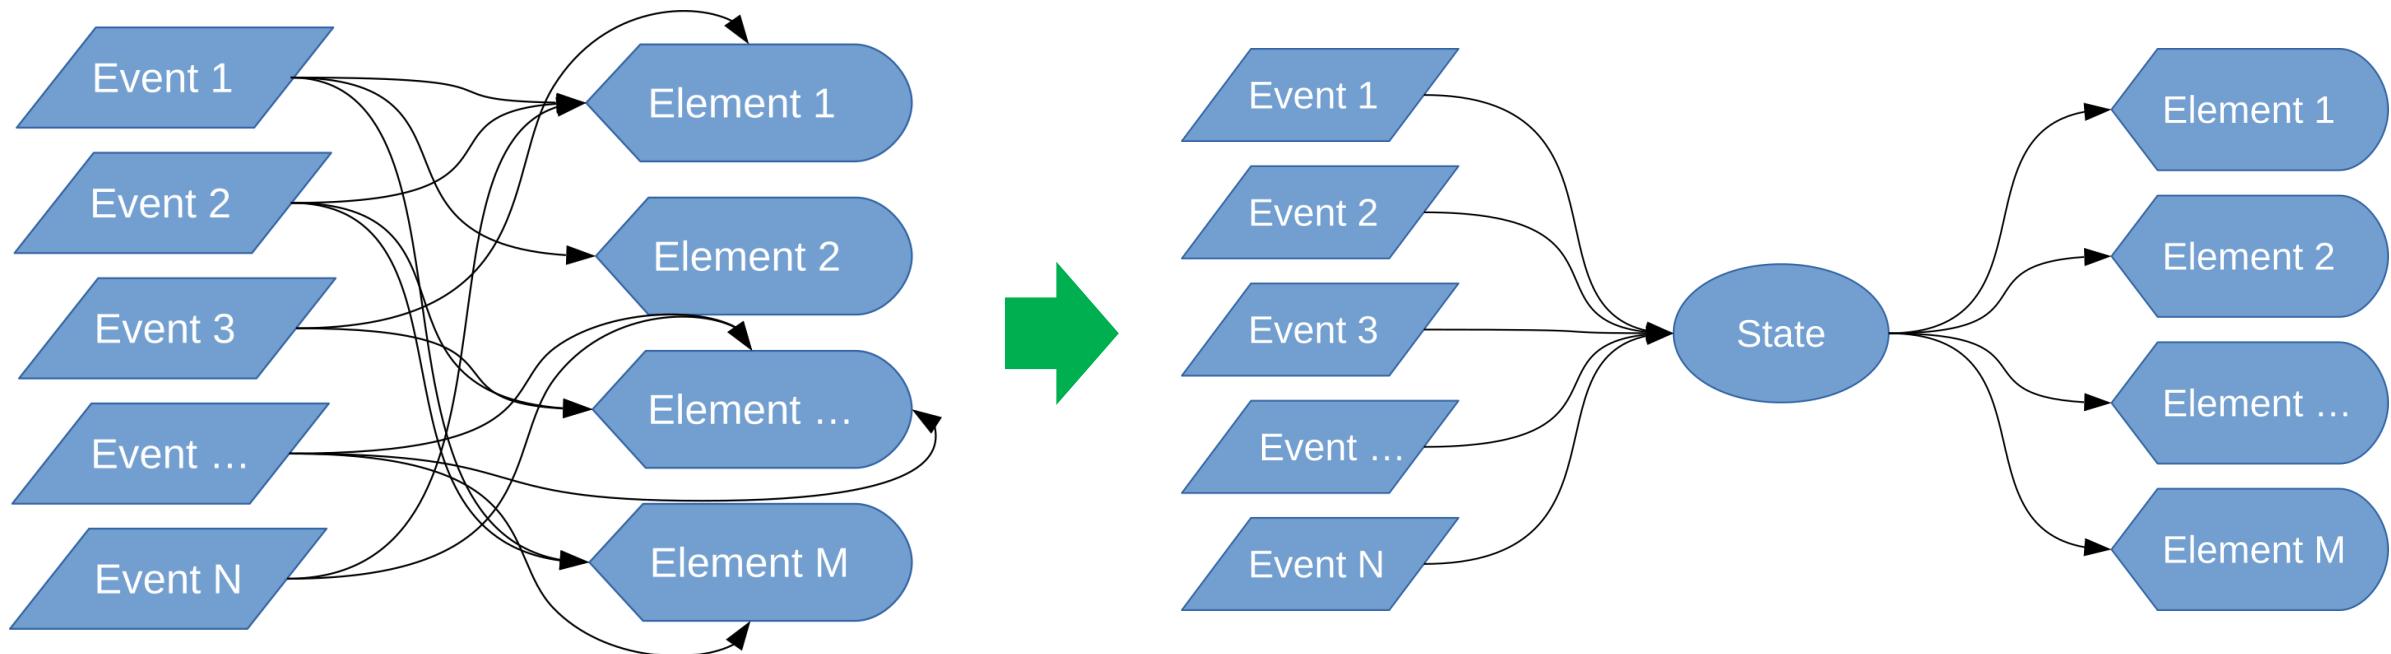
For developers, the question is whether they want to learn yet **one more layer of abstraction** to their coding. On one hand, there's the elusive promise of a **common API to streaming engines** that in theory should let you mix and match, or swap in and swap out.

*Tony Baer (Ovum) @ ZDNet - Apache Beam and Spark:
New competition for squashing the Lambda Architecture?*

Direct Acyclic Graphs - DAG

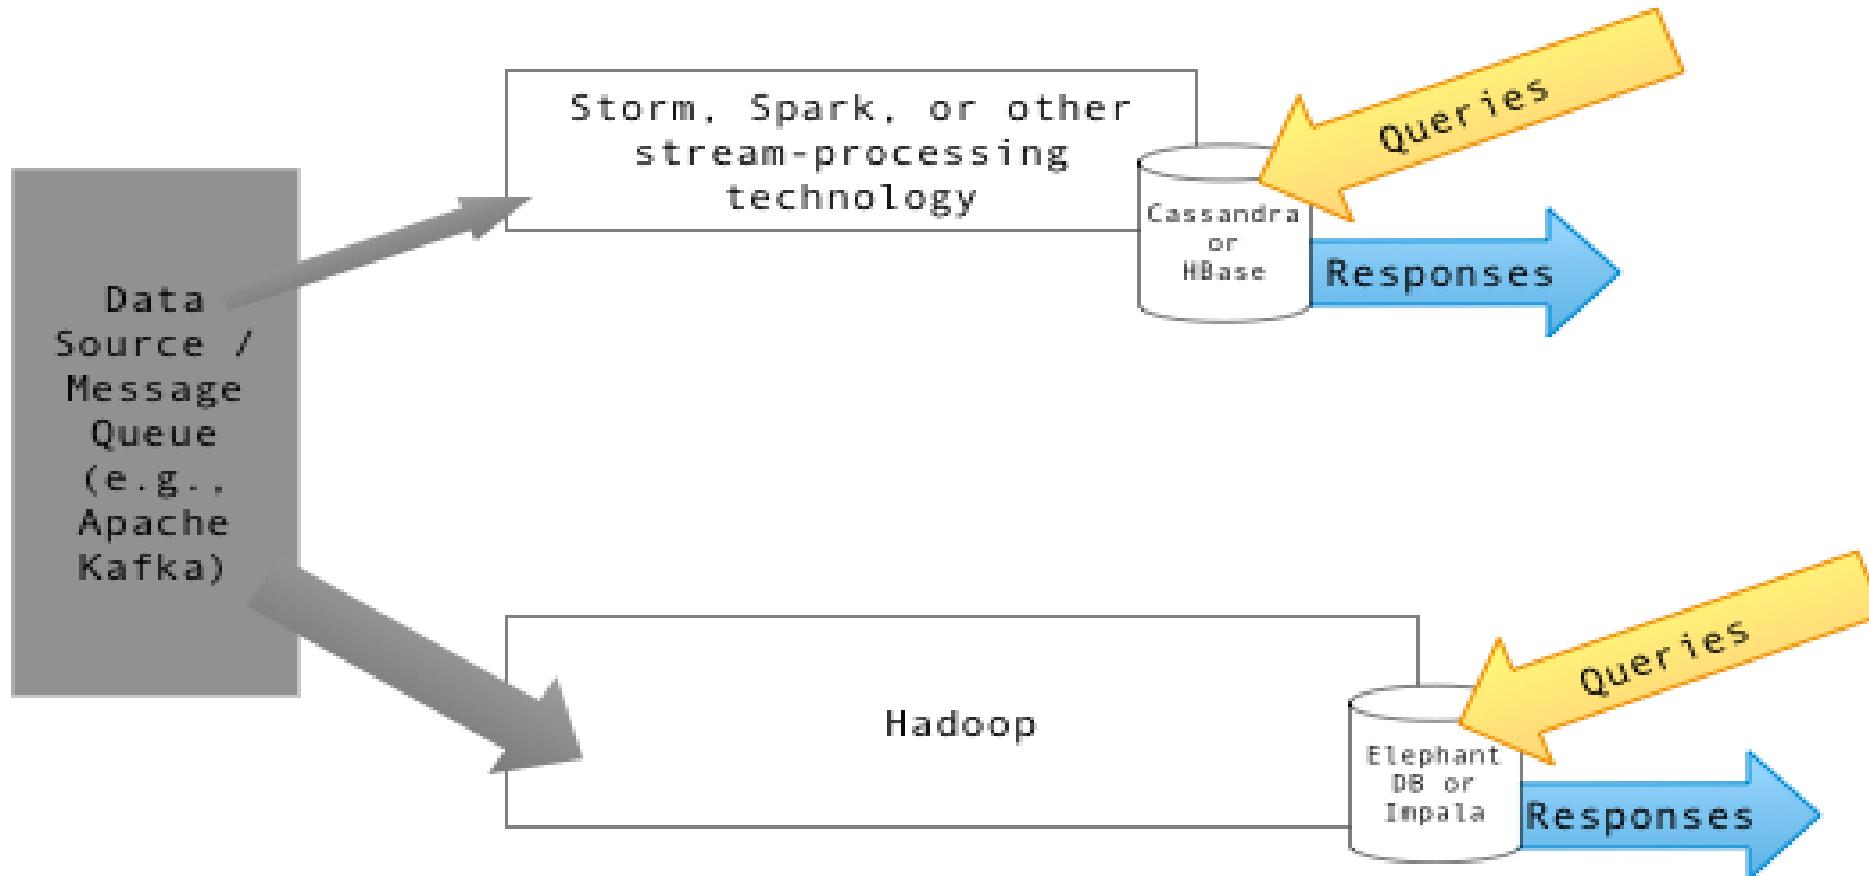


Event Sourcing – Events vs. State (Snapshots)



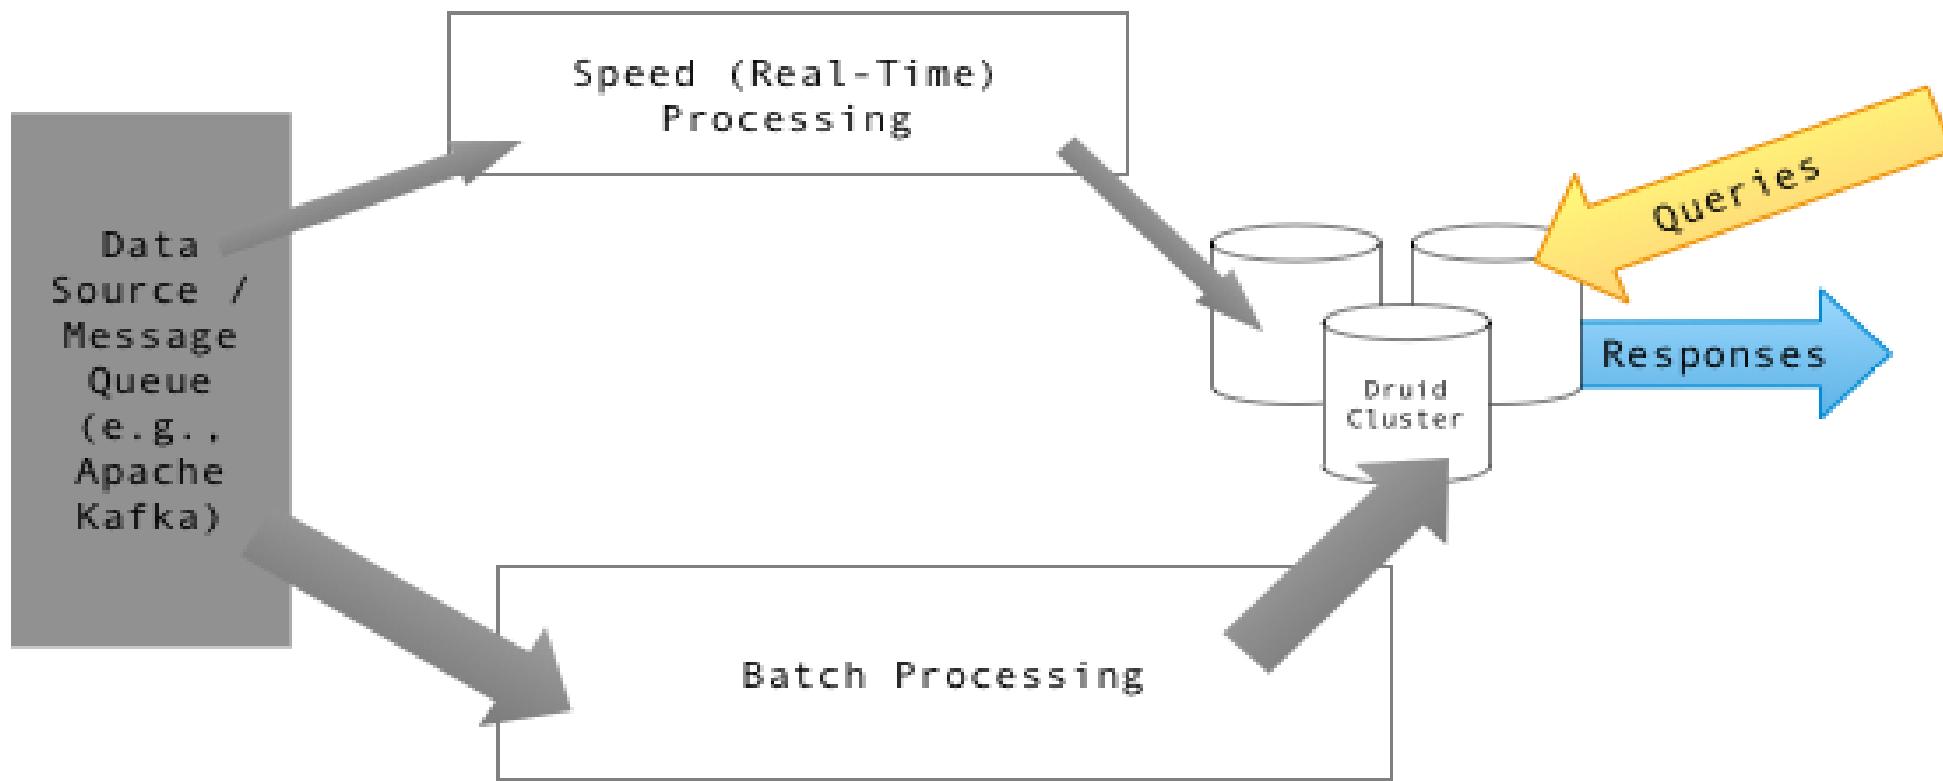
Lambda Architecture - I

Query = λ (Complete data) = λ (live streaming data) * λ (Stored data)

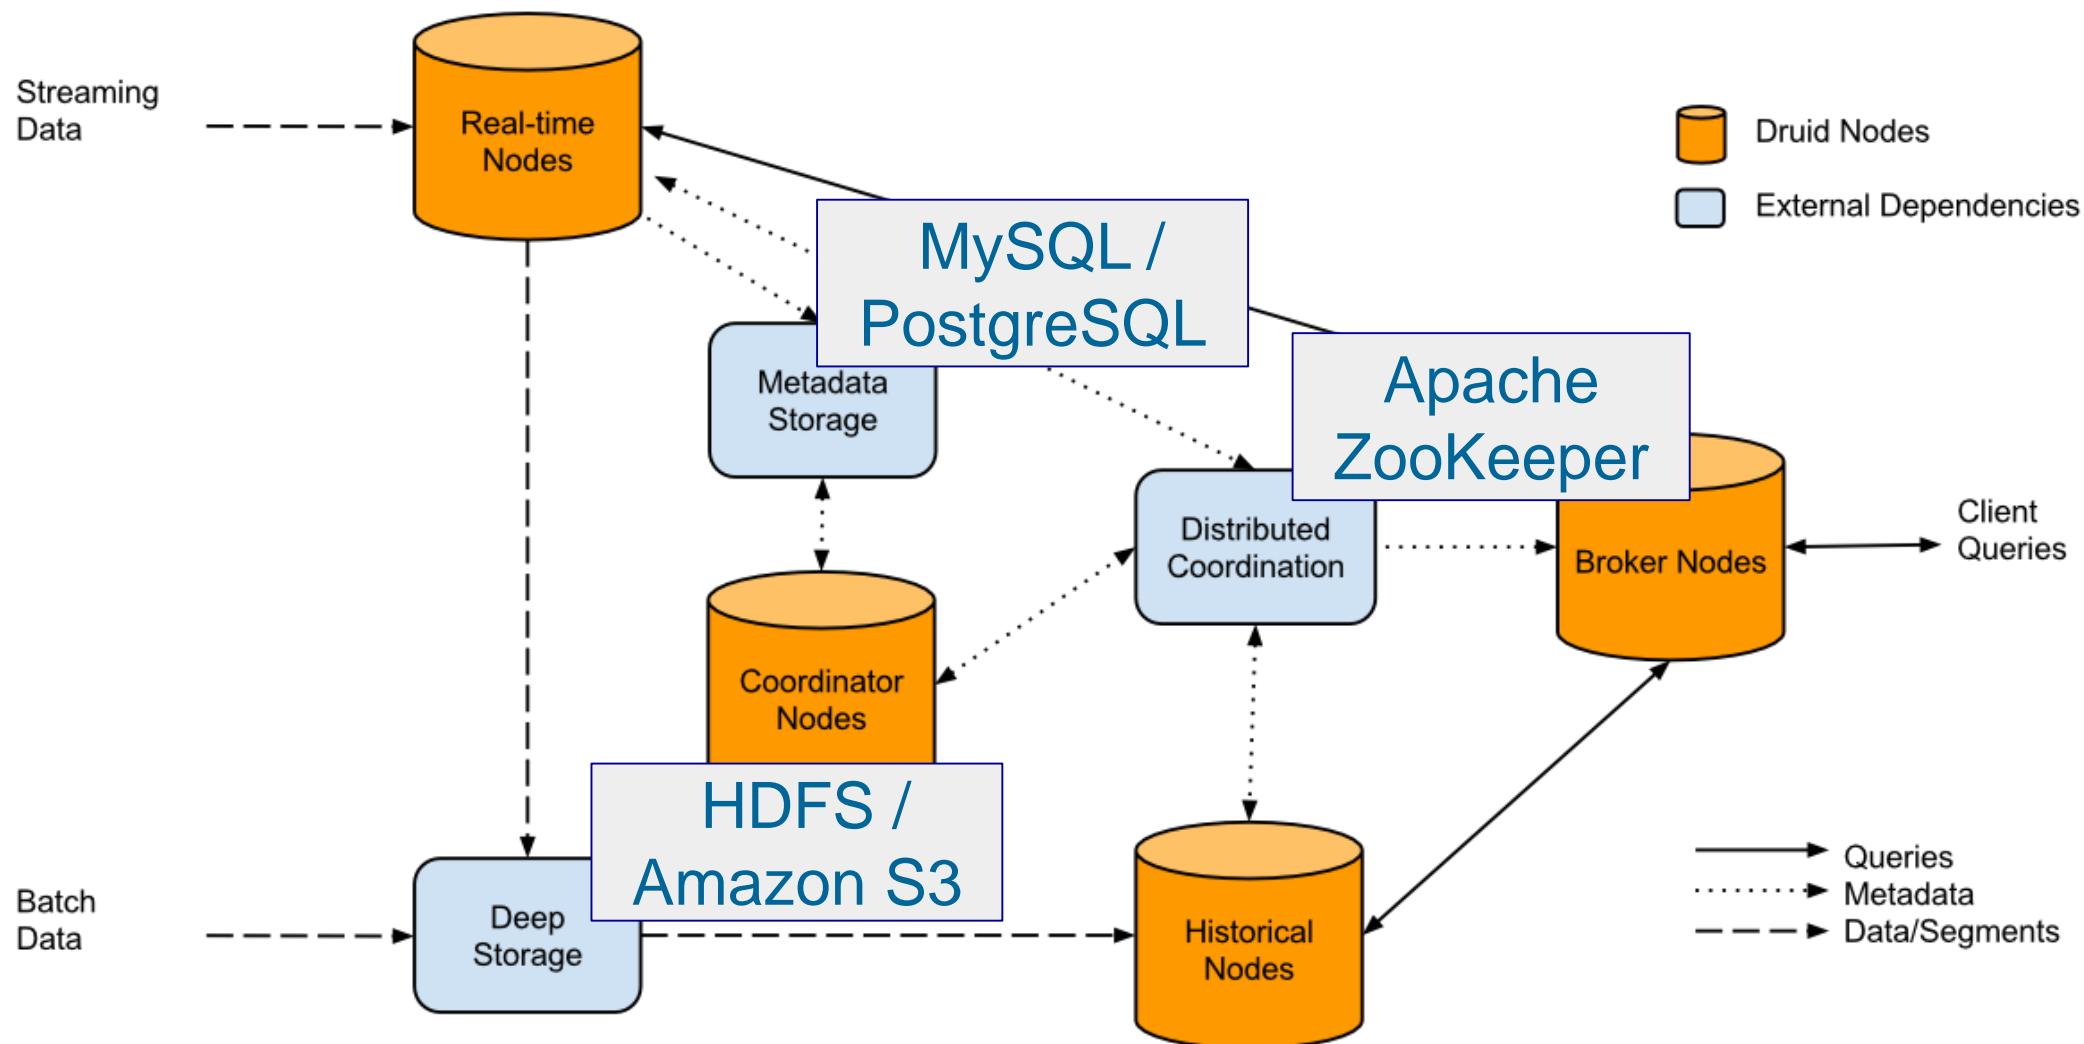


Lambda Architecture - II

Query = λ (Complete data) = λ (live streaming data) * λ (Stored data)



Lambda Architecture - Druid Distributed Data Store



Kappa Architecture

Query = K (New Data) = K (Live streaming data)

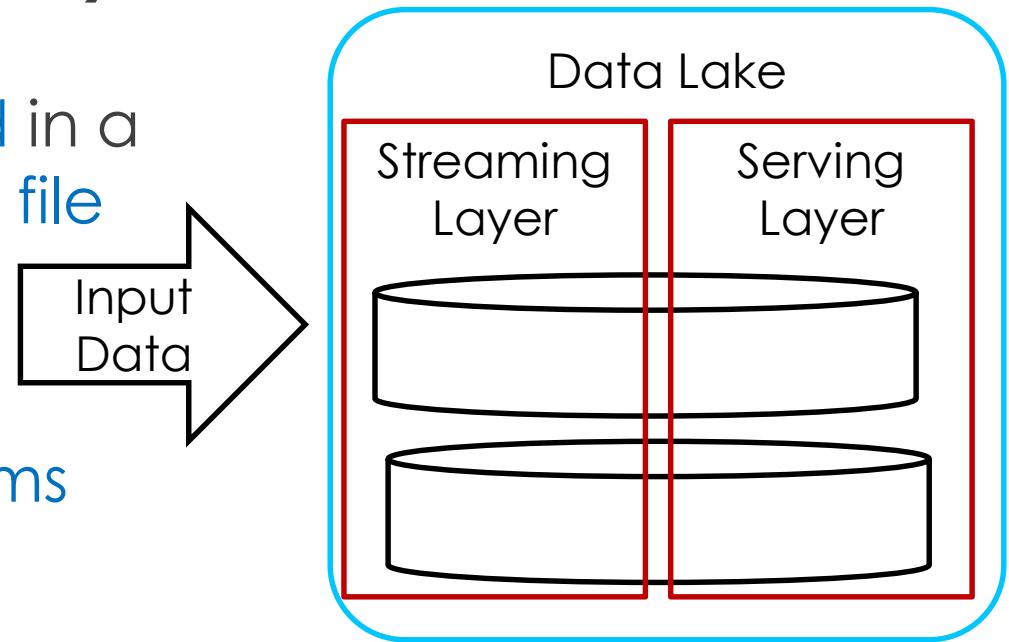
- Proposed by Jay Kreps in 2014
- Real-time processing of distinct events
- Drawbacks of Lambda architecture:
 - It can result in coding overhead due to comprehensive processing
 - Re-processes every batch cycle which may not be always beneficial
 - Lambda architecture modeled data can be difficult to migrate
- Canonical data store in a Kappa Architecture system is an append-only immutable log (like Kafka, Pulsar)



Kappa Architecture II

Query = K (New Data) = K (Live streaming data)

- Multiple **data events or queries** are logged in a queue to be catered against a **distributed file system storage** or **history**.
- The order of the events and queries is not predetermined. **Stream processing platforms** can interact with **database** at any time.
- It is **resilient** and **highly available** as handling **terabytes of storage** is required for each node of the system to **support replication**.
- Machine learning is done on the **real time basis**



Zeta Architecture

- Main characteristics of Zeta architecture:
 - file system ([HDFS](#), [S3](#), [GoogleFS](#)),
 - realtime data storage ([HBase](#), [Spanner](#), [BigTable](#)),
 - modular processing model and platform ([MapReduce](#), [Spark](#), [Drill](#), [BigQuery](#)),
 - containerization and deployment ([cgroups](#), [Docker](#), [Kubernetes](#)),
 - Software solution architecture ([serverless computing](#) – e.g. [Amazon Lambda](#))
- [Recommender systems](#) and [machine learning](#)
- Business applications and dynamic global resource management ([Mesos + Myriad](#), [YARN](#), [Diego](#), [Borg](#)).

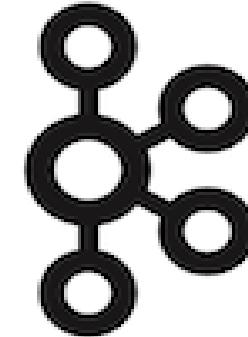
Distributed Stream Processing – Apache Projects:

- Apache Spark is an open-source cluster-computing framework. Spark Streaming, Spark Mllib
- Apache Storm is a distributed stream processing – streams DAG
- Apache Samza is a distributed real-time stream processing framework.

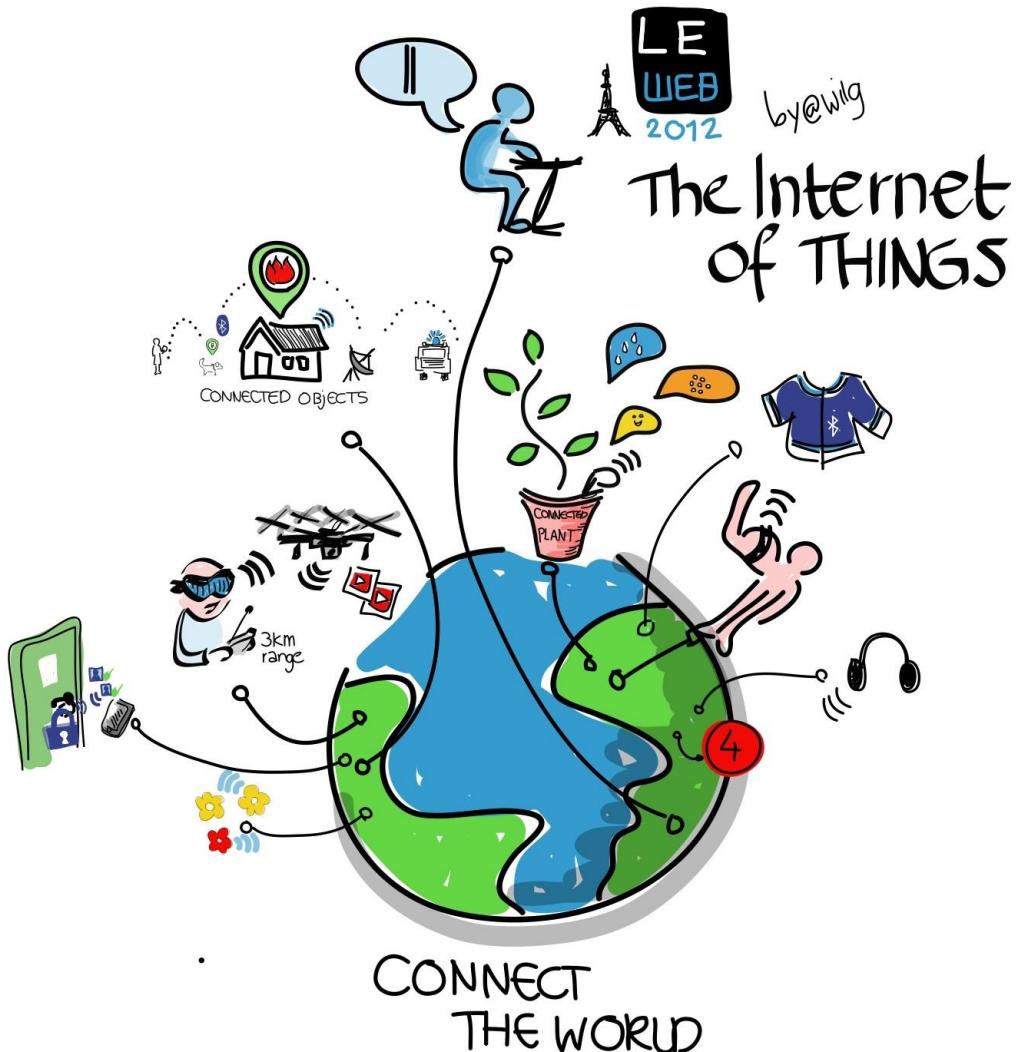


Distributed Stream Processing – Apache Projects II

- [Apache Flink](#) - open source stream processing framework – Java, Scala
- [Apache Kafka](#) - open-source stream processing (Kafka Streams), real-time, low-latency, high-throughput, massively scalable pub/sub
- [Apache Beam](#) – unified batch and streaming, portable, extensible



Example: Internet of Things (IoT)



Wearable Electronics :)

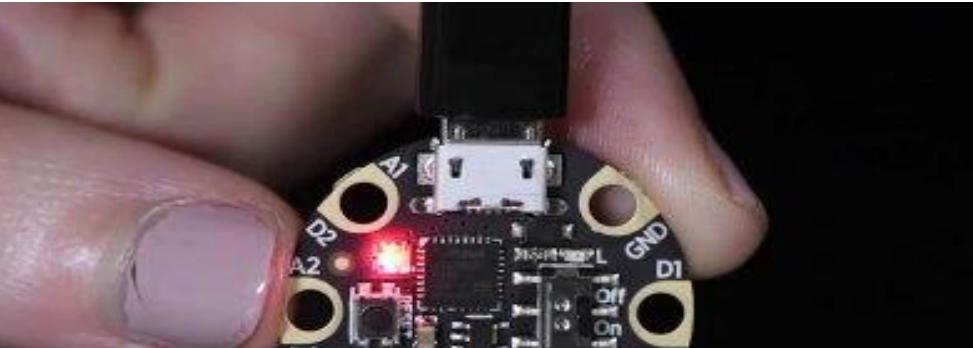
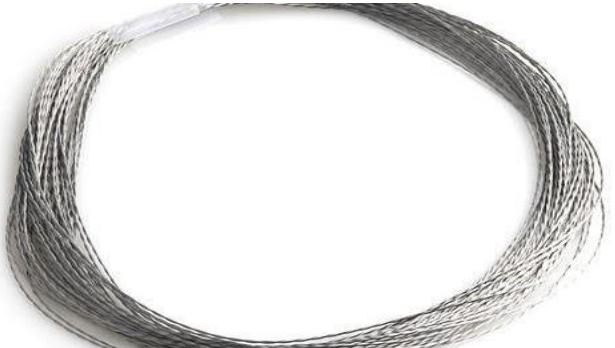
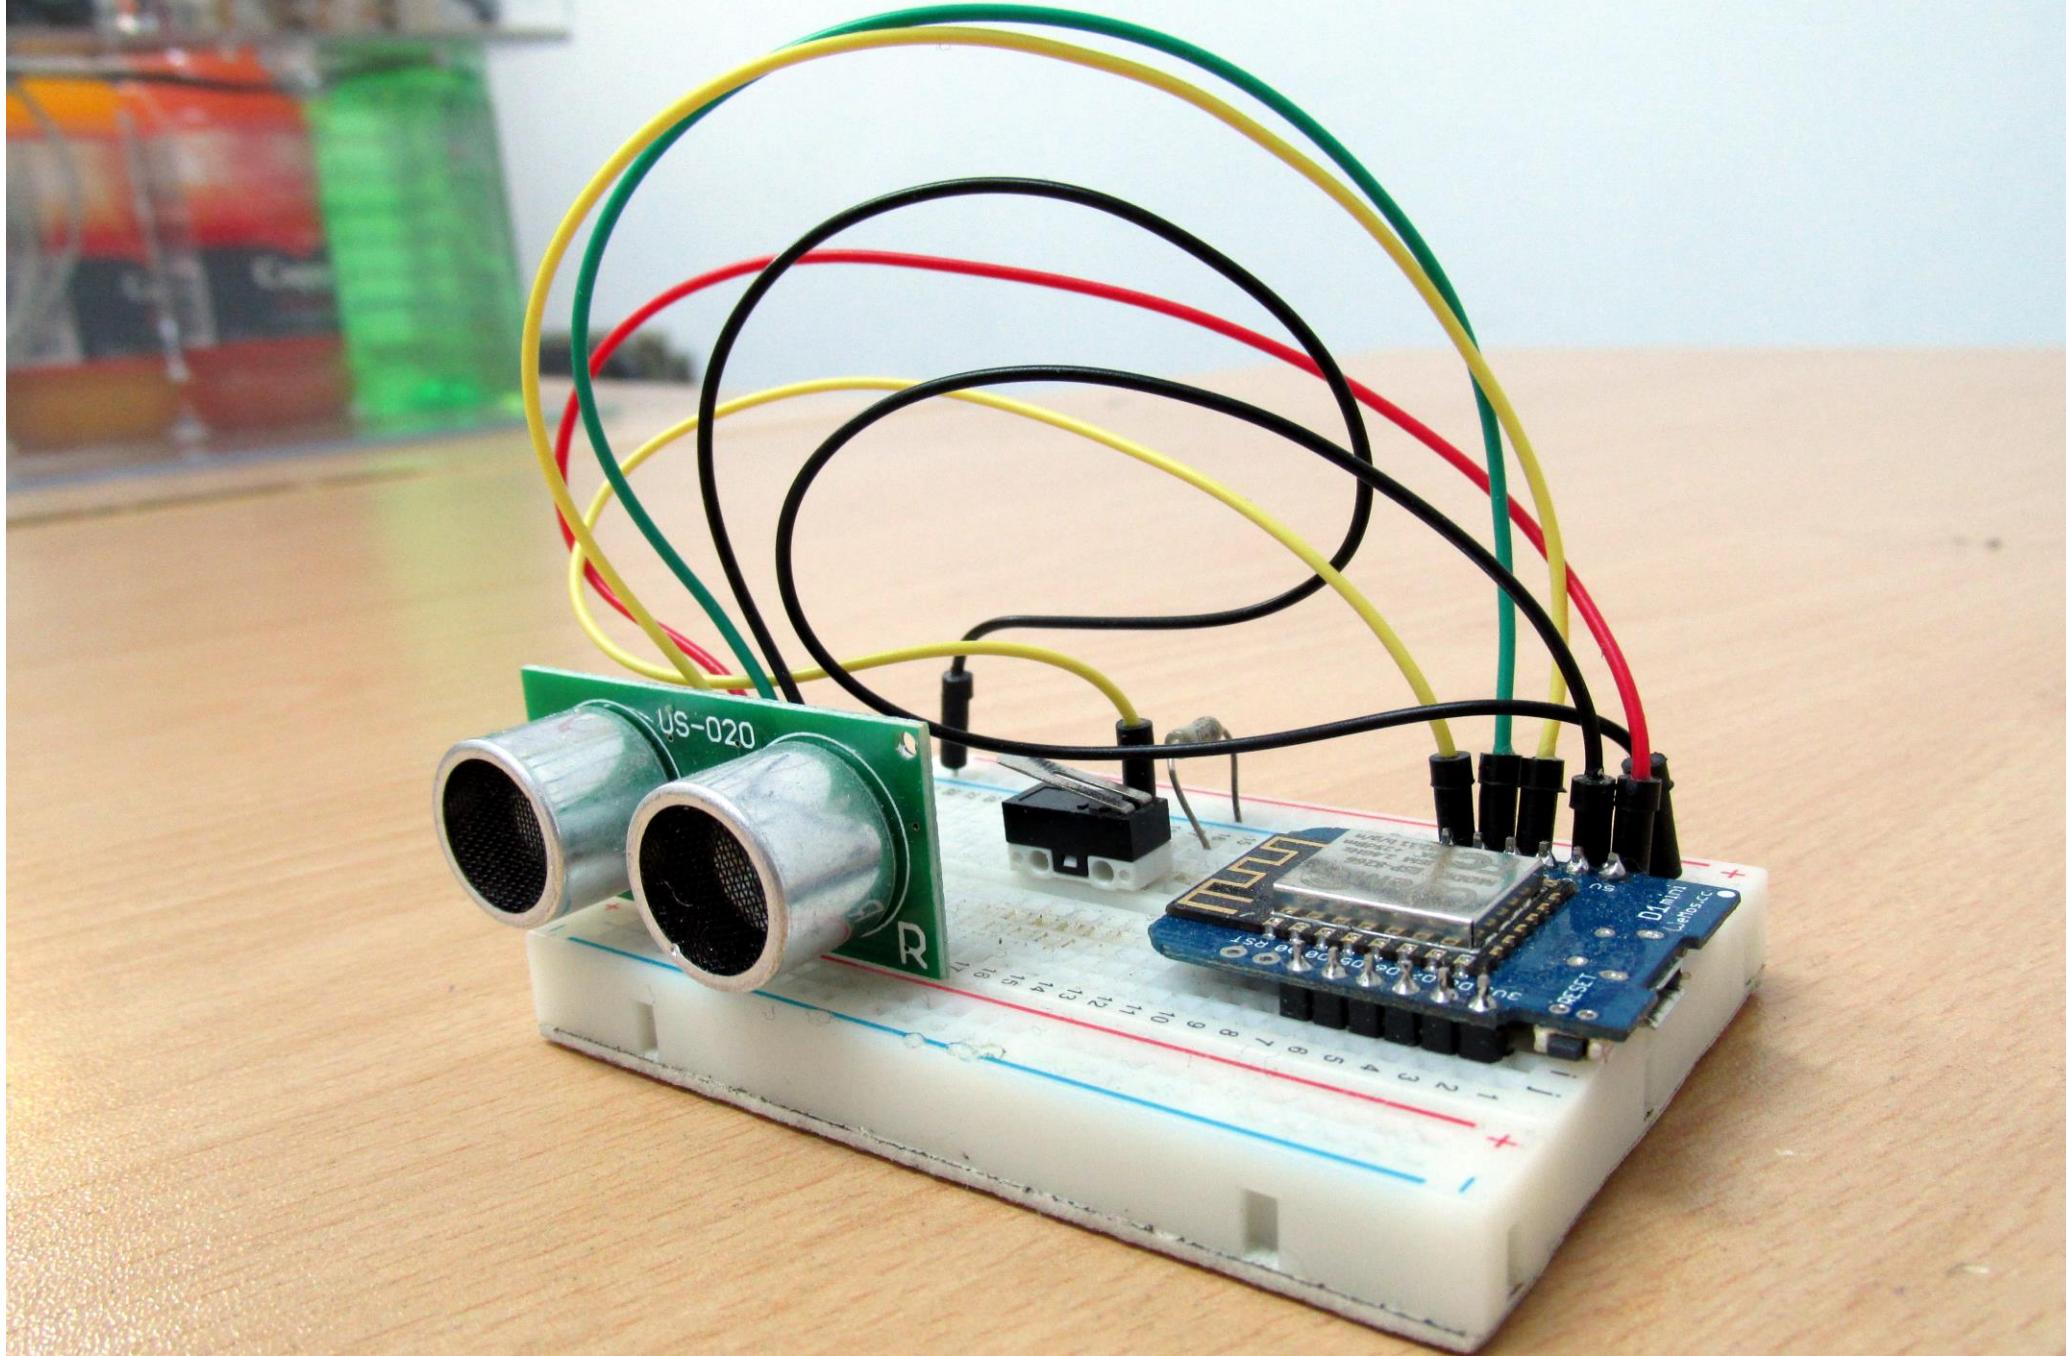
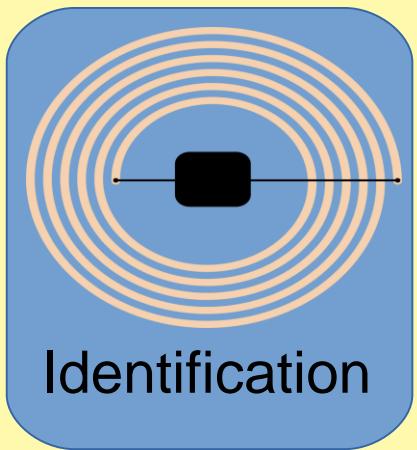


Photo credit: Adafruit, Becky Stern and others



Key Elements of IoT

Internet of Things (IoT)



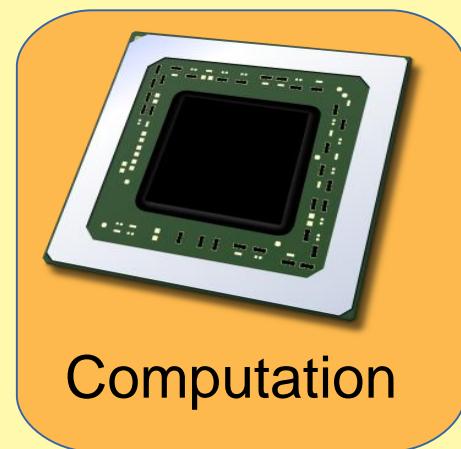
Identification



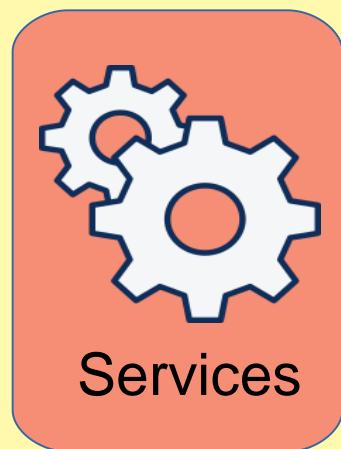
Sensors



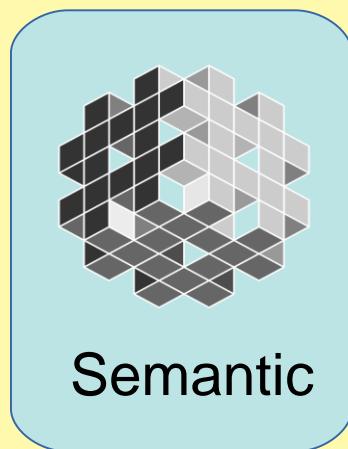
Connectivity



Computation

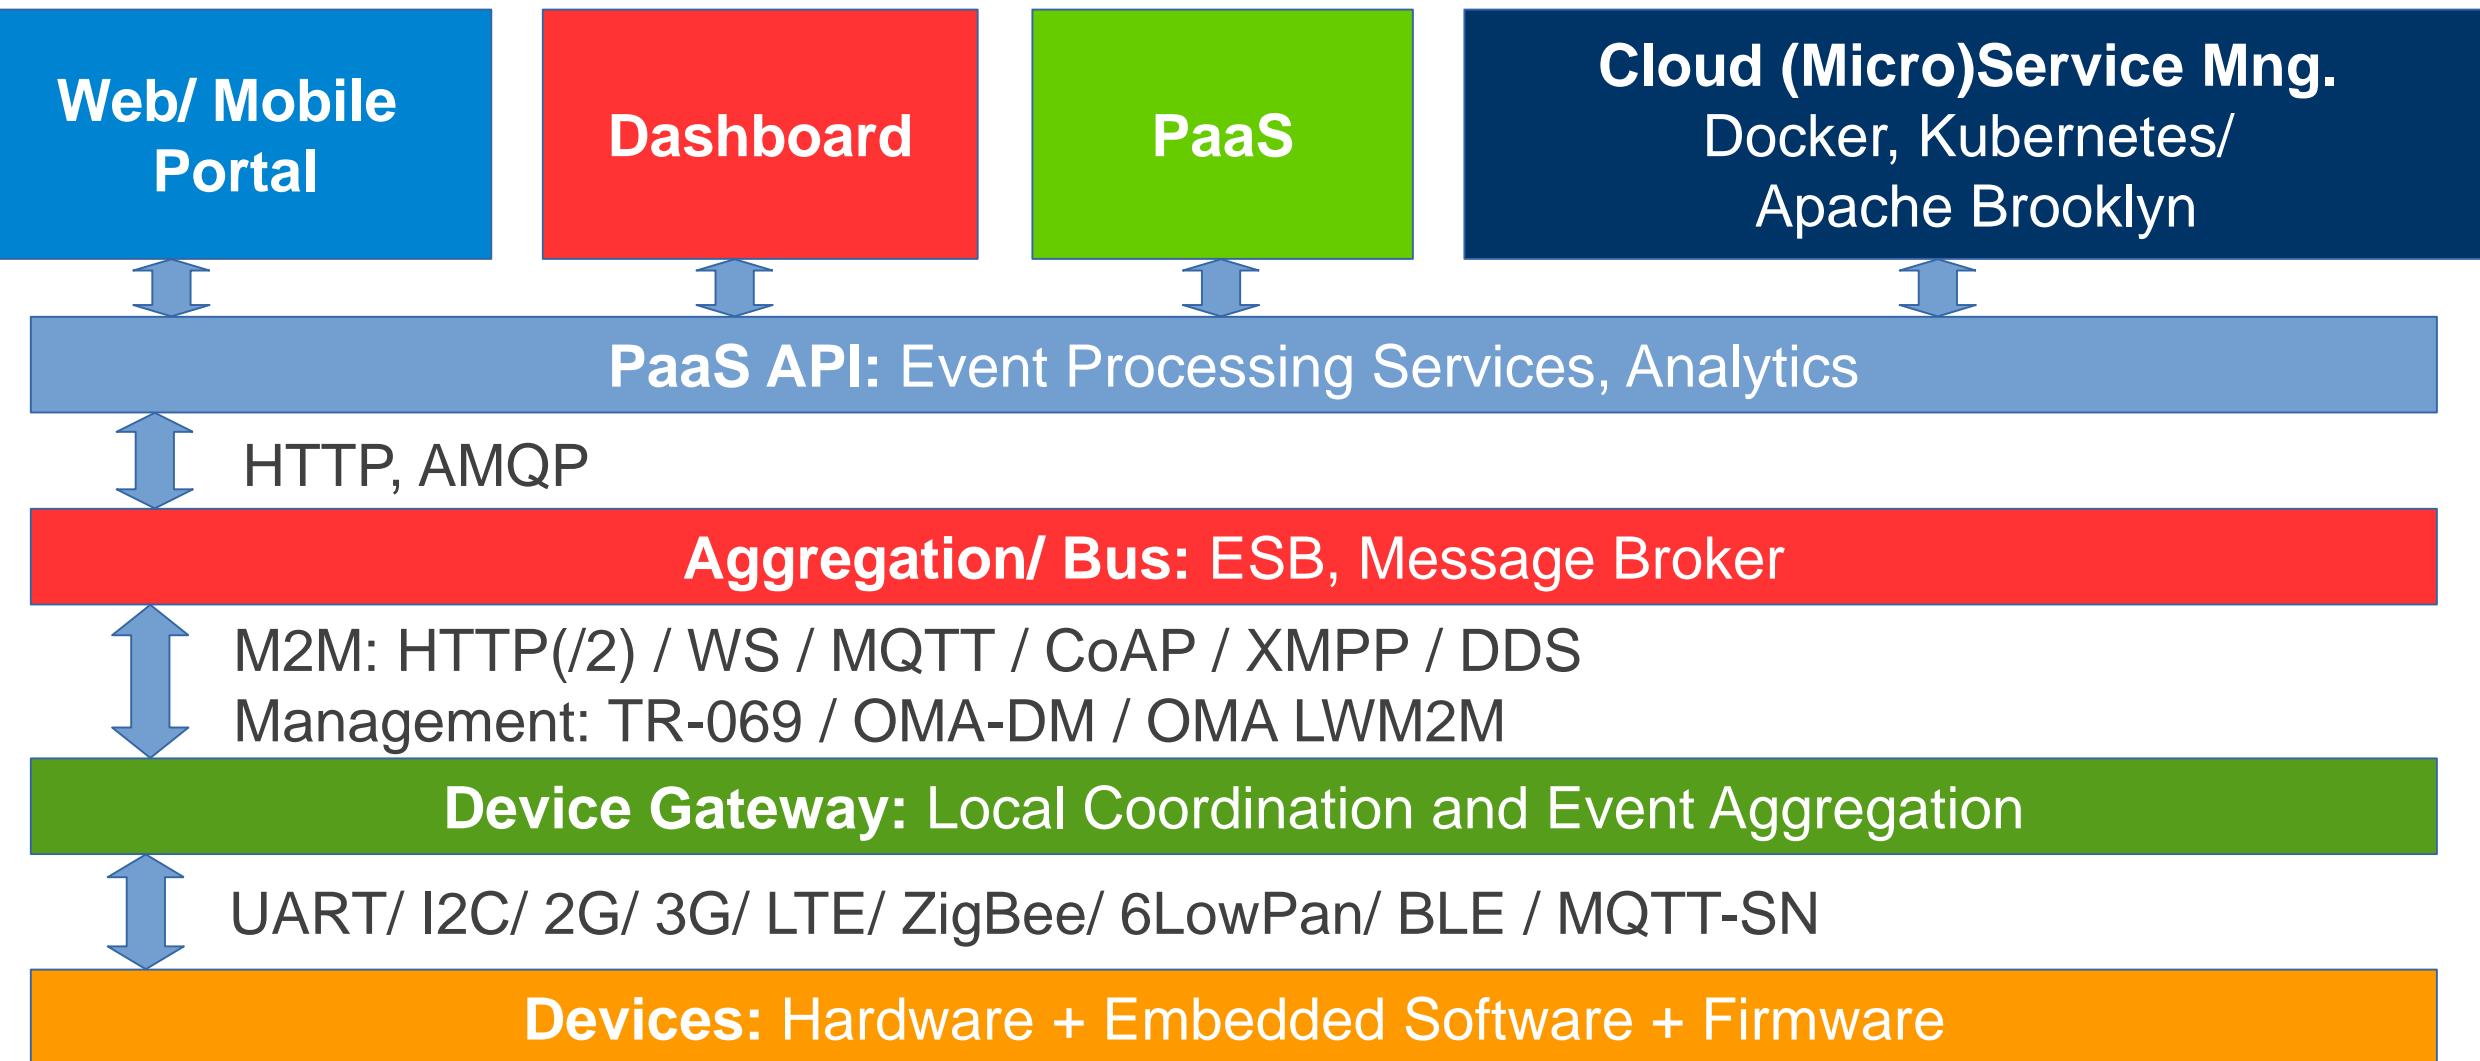


Services



Semantic

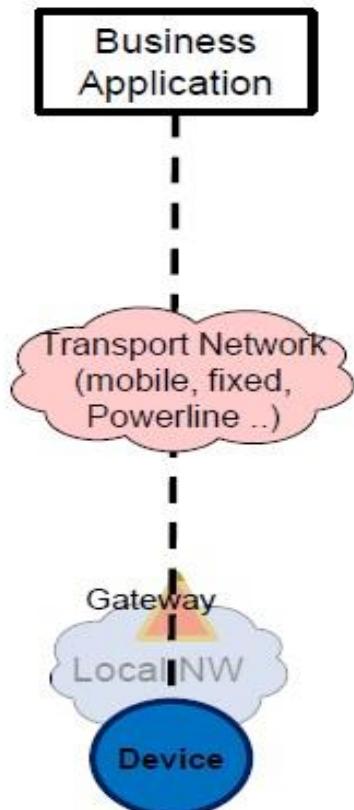
IoT Services Architecture



Vertical vs. Horizontal IoT

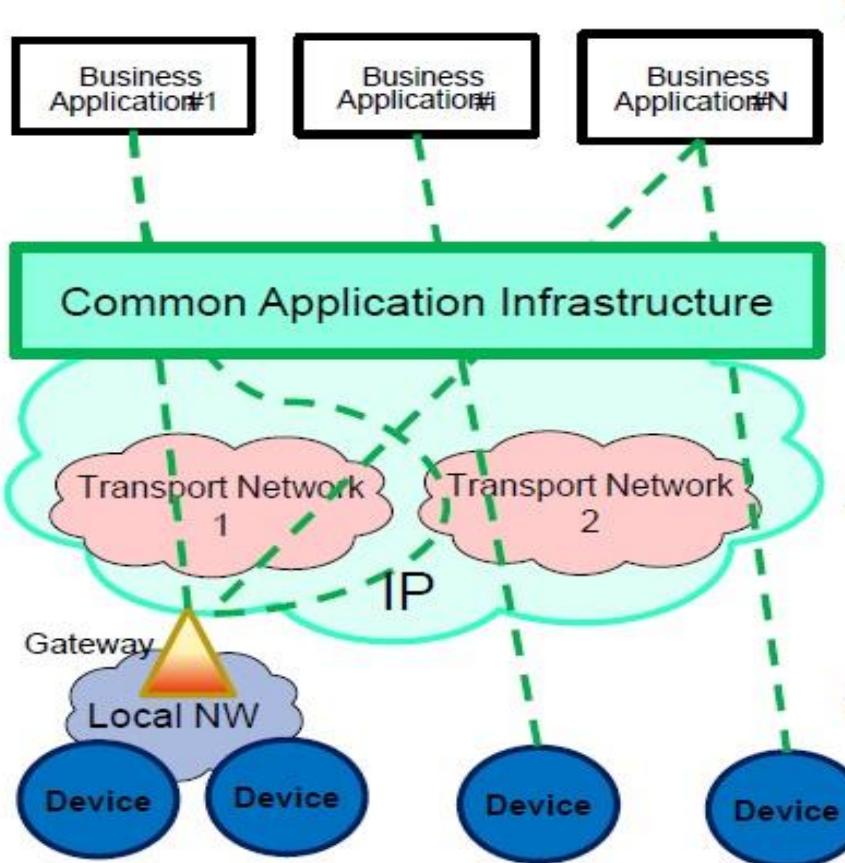
Pipe (vertical):

1 Application, 1 NW,
1 (or few) type of Device



Horizontal (based on common Layer)

Applications share common infrastructure, environments
and network elements



M2M Applications providers run individual M2M services. Customer is Device owner

M2M Service provider hosts several M2M Applications on his Platform.

Wide Area Transport Network operator(s) Customer is the M2M service provider

End user owns / operates the Device or Gateway

Cloud, Fog and Mist Computing

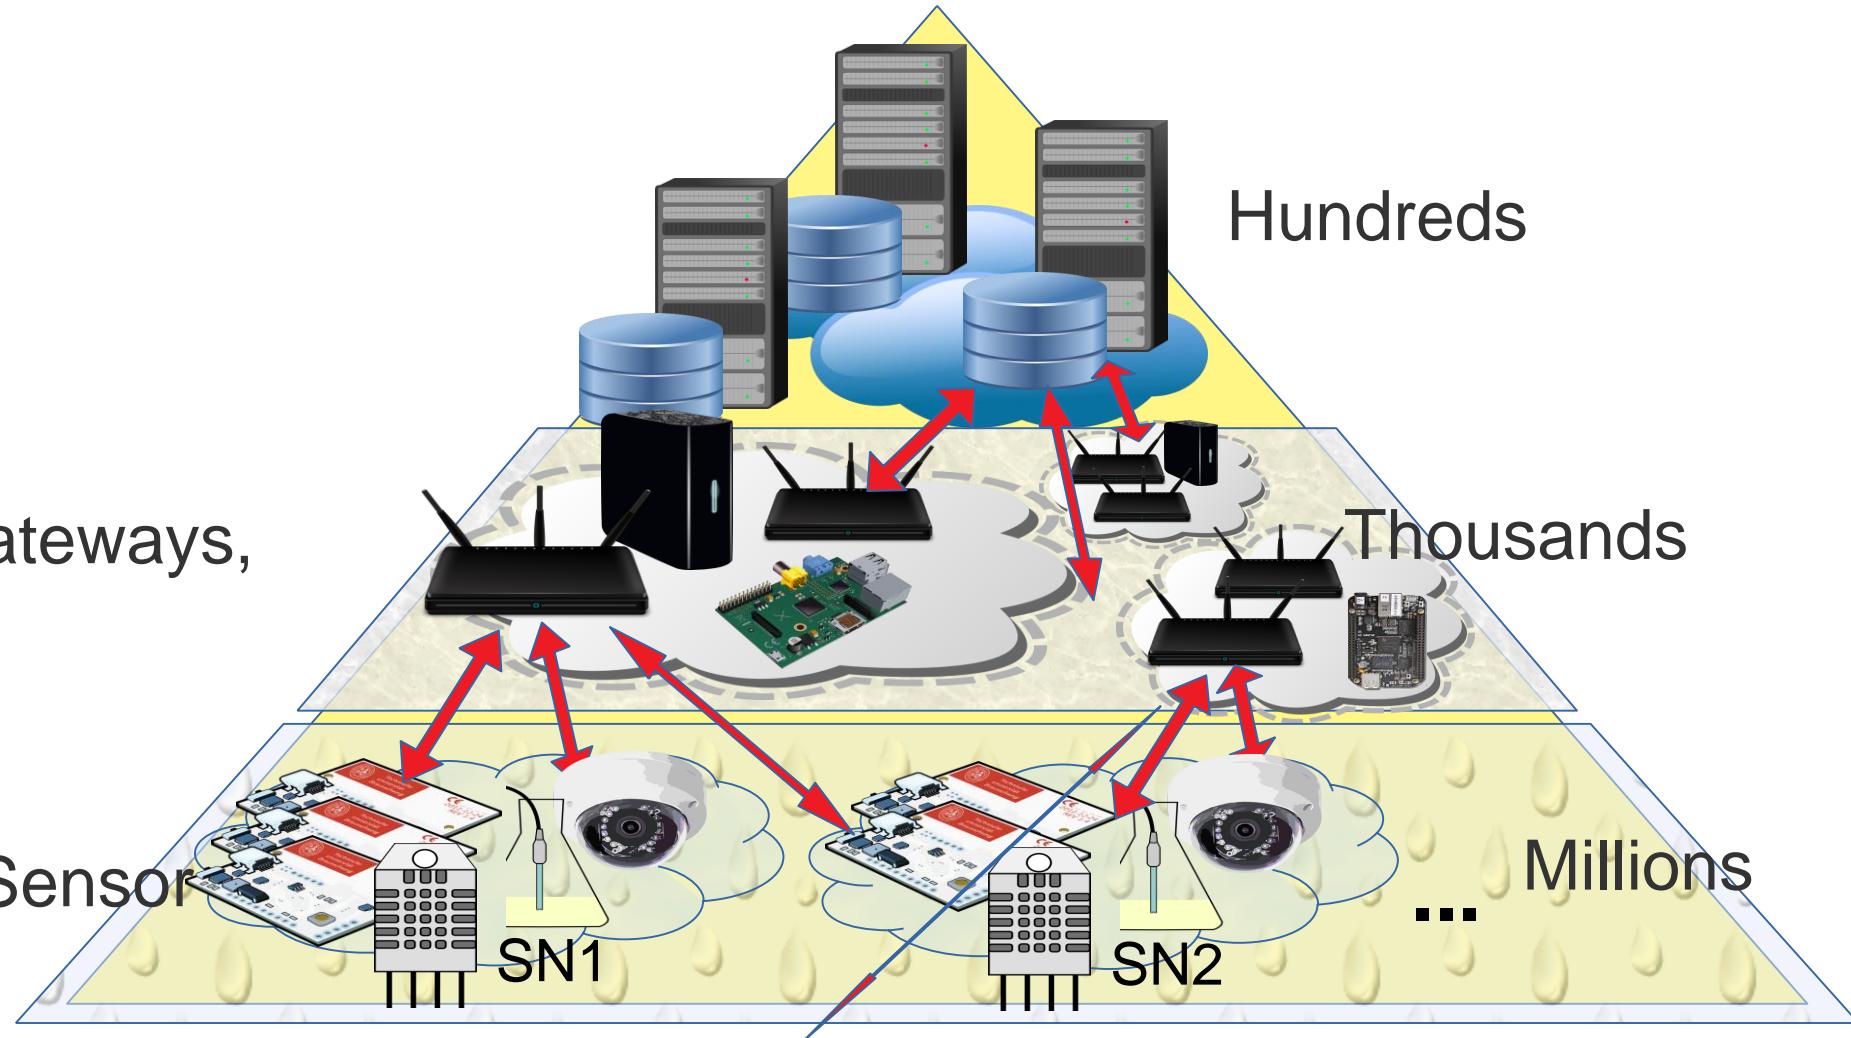
Cloud Computing
(Data-centers)



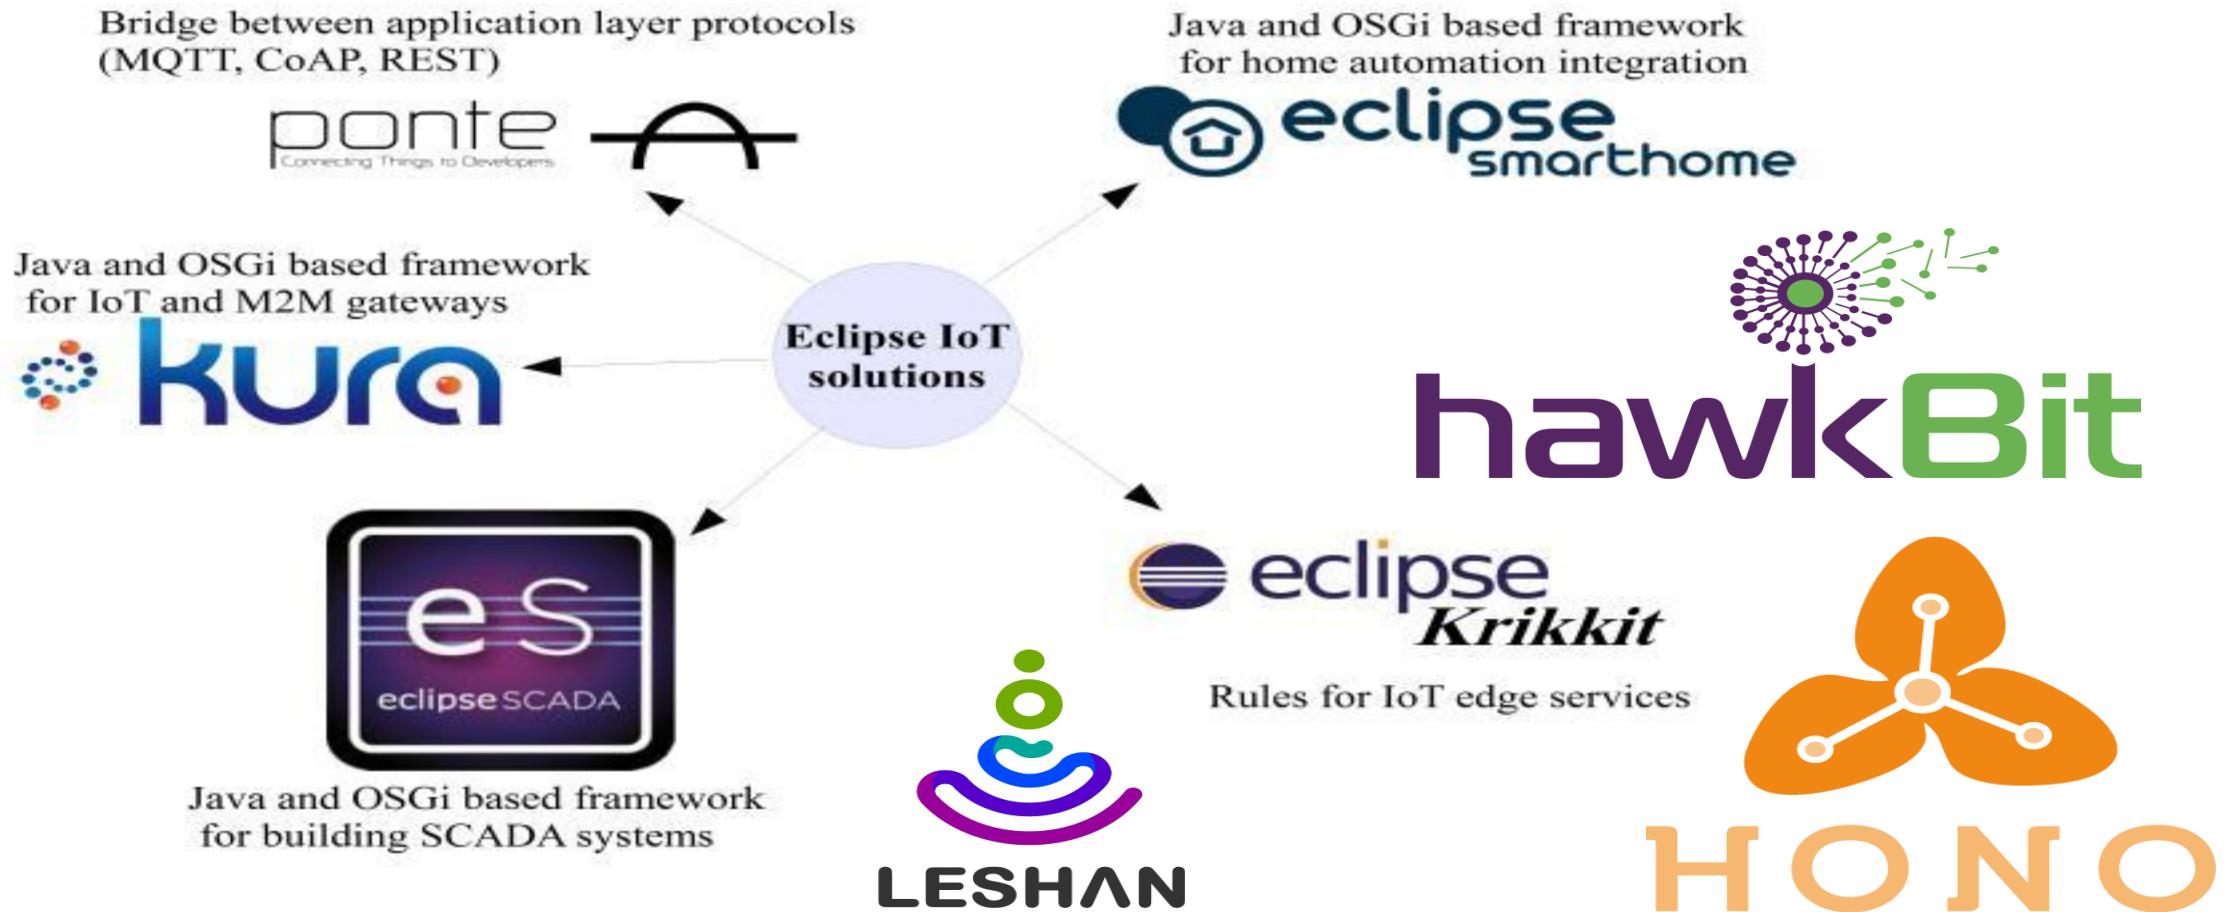
Fog Computing
(Fog Nodes, Edge Gateways,
Cloudlets)



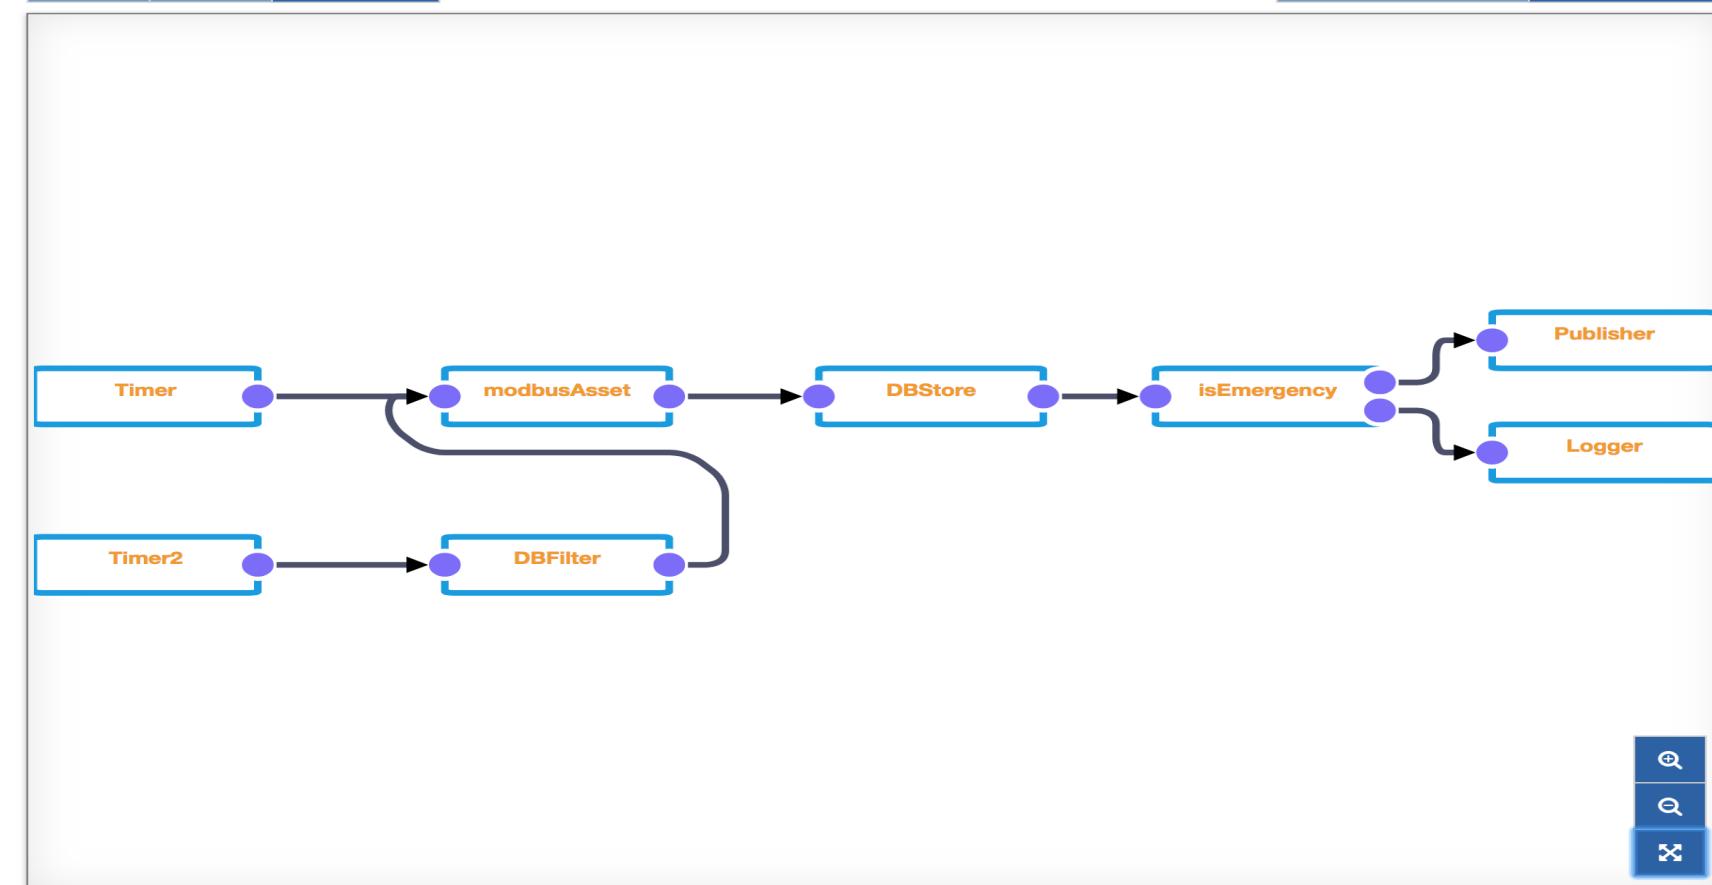
Mist Computing
(Smart IoT Devices, Sensor
Networks)



Eclipse IoT Platform



Wire Graph

[Apply](#) [Reset](#) [Download](#)[Delete Component](#) [Delete Graph](#)

Wire Components

- Subscriber
- Timer
- ← Logger
- ← Publisher
- ↔ RegexFilter
- ↔ Asset
- ↔ Fifo
- ↔ HSQL DB Filter (Deprecated)
- ↔ H2 DB Store
- ↔ HSQL DB Store (Deprecated)
- ↔ H2 DB Filter
- ↔ Conditional
- ↔ Join

System

- ! Status
- Device
- Network
- Firewall
- Cloud Services
- Drivers and Assets

Wires

Packages

Settings

Services

 Search +

Simple Artemis MQTT Broker

ActiveMQ Artemis Broker

ClockService

DeploymentService

>_ CommandService

WebConsole

H2DbService

PositionService

RestService

File Alarms Window Help

Eclipse SCADA Demo Client

2014-08-04 10:24:56 es

Main

Main
 Borders
 Shapes

demo.openscada.org scada.eclipse.org Eclipse Greenhouse Sim Greenhouse

CPU Idle 99 % CPU Idle 95 % Temperature 24.9 °C Temperature 26.0 °C

Load Average 0.1 Load Average 0.3

Arduino fortiss Living Lab

Light 168.5 PV DC 2876
 Temperature 28.0 °C PV AC 2761
 Humidity 44.0 %

Arduino

Temperature 26 01:53 100%
 Humidity 46.97 01:53 100%
 Light 0 01:53 100%

Time 03:00 08:00 13:00 18:00 23

6 Summary

Overview Item List

Measured Data

168.5 Light
 28.0 Temperature (°C)
 44.0 Humidity (%)

Webcam

2014-08-04 16:24:56-01

Control

Light On
 Light Off

admin @ Sample Context no Alarm

Eclipse Smart Home – Open HAB

03/09/2015

12:47 PM

Front Porch
Light



Outdoor
Lights



Away Mode



Garage Door



Indoor
Temperature
23.7 °C

Today

15 °C

16° / 9° ☂ 40% ☀ 68%

Tomorrow

17° / 10° ☂ 10%

Living Room
Lights



Upstairs Landing



0 %



Front Room
Lights



Hallway



0 %



03-09-2015 · Thu



CCTV
Cameras

Room
Lighting



01/09/2017

1:40 AM



57.4 °F

Humidity: 52%
Pressure: 30.12 in

Today



Mostly Cloudy

75 °F
46 °F

Tuesday



Clear
70 °F
45 °F

Wednesday



Clear
66 °F
45 °F

Utility



Hallway



Family Room



Living Room



Living Room Lights

Back Yard



73.0 °F

Office: 77°

Front Yard



Porch



Driveway

Kitchen



Stereo Volume

17%

Stereo

70

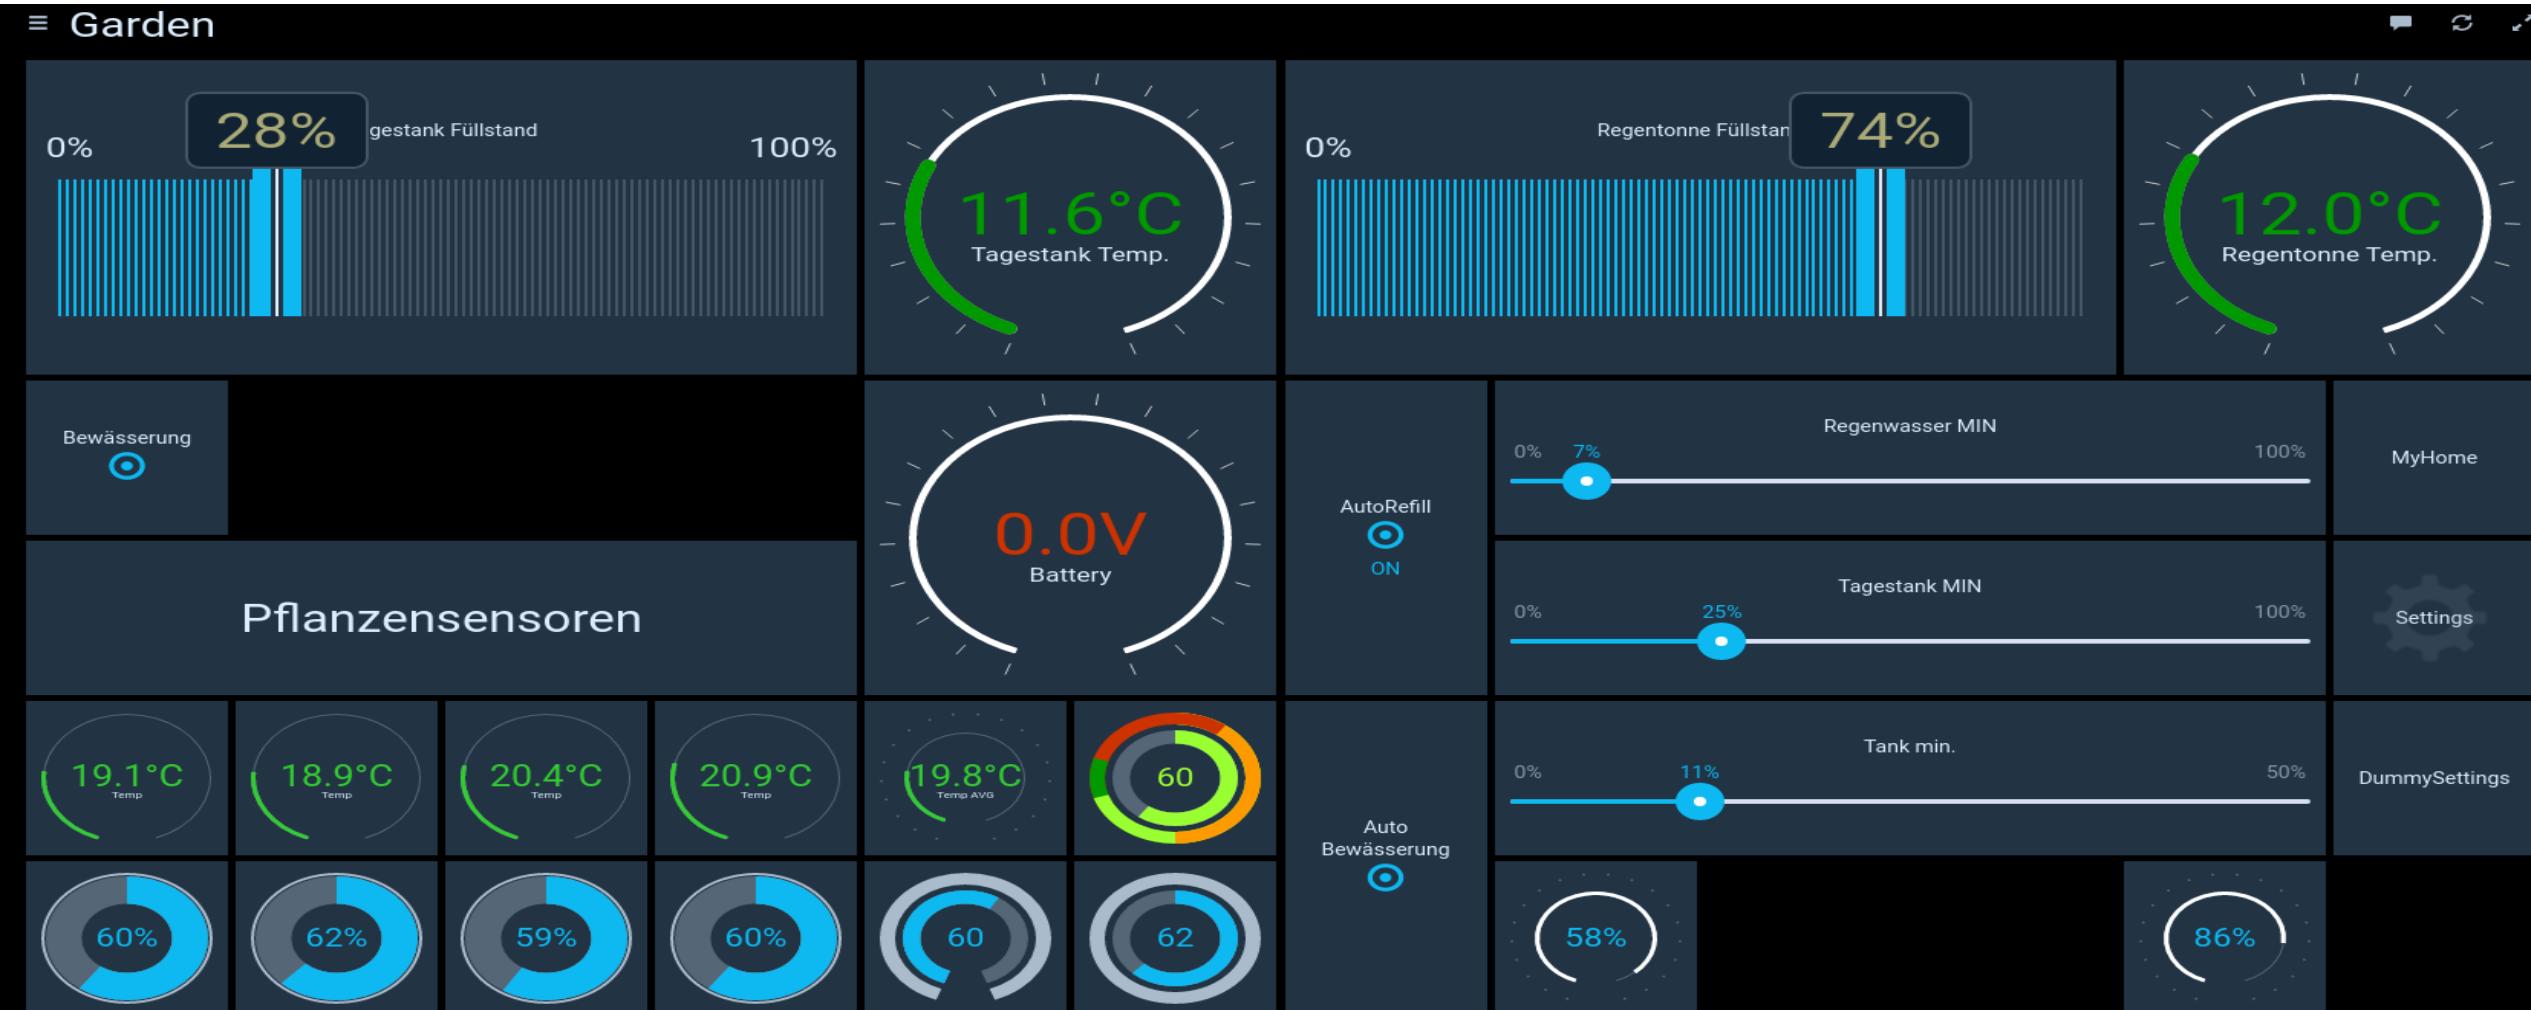


34%

Office: 77°

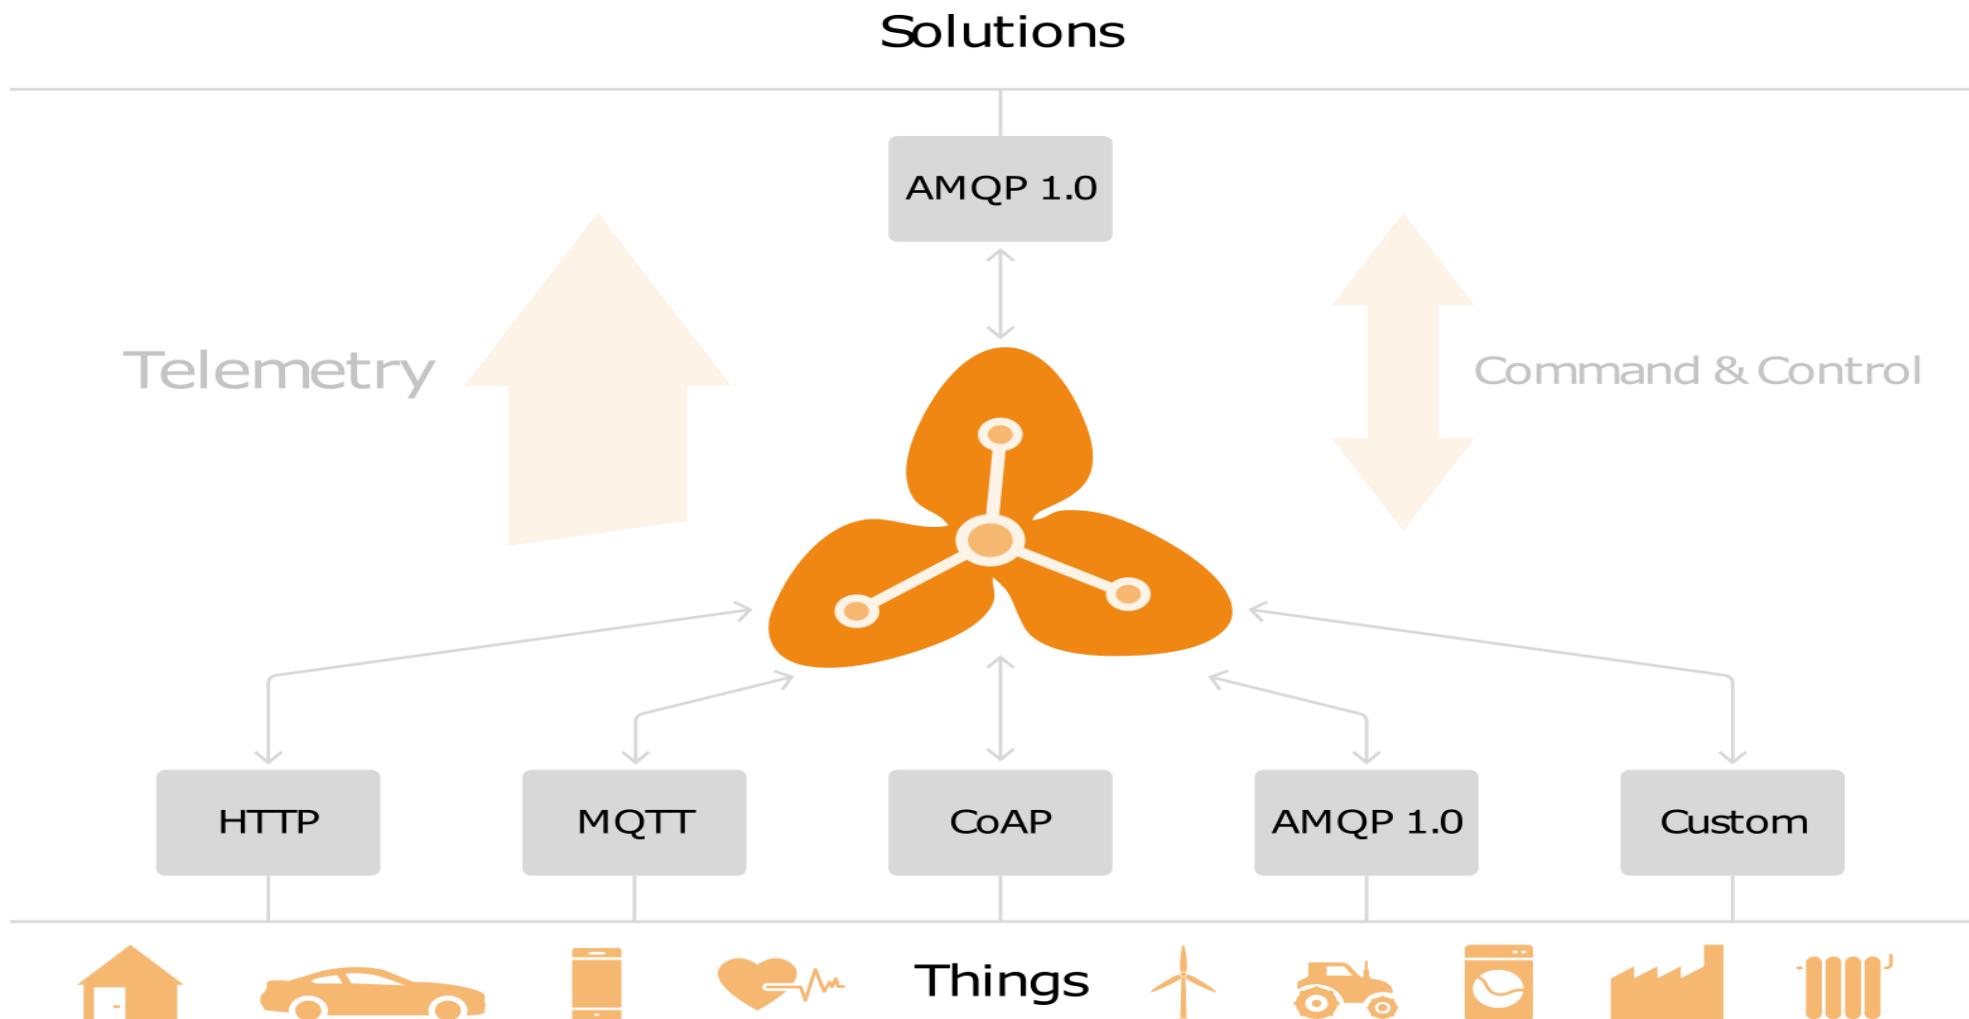
Wake Up

Open HAB: Hydroponics Dashboard



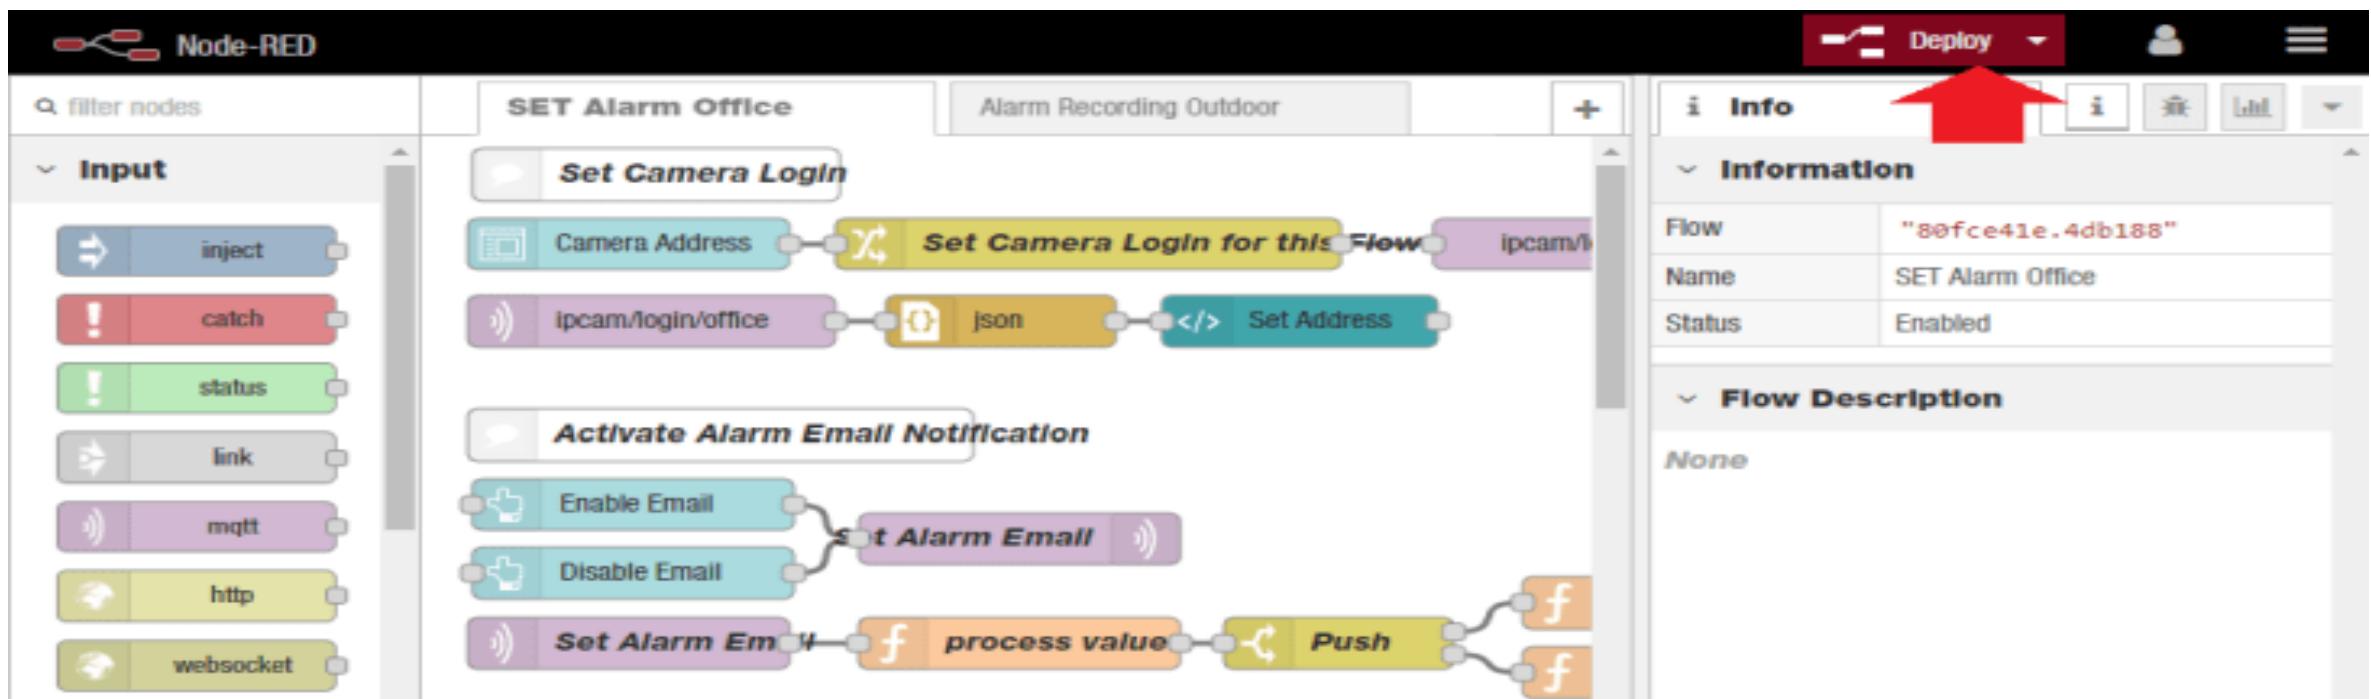
Source: <https://community.openhab.org/t/aquaponics-goes-openhab/26456>

Eclipse Hono: Connect. Command. Control.

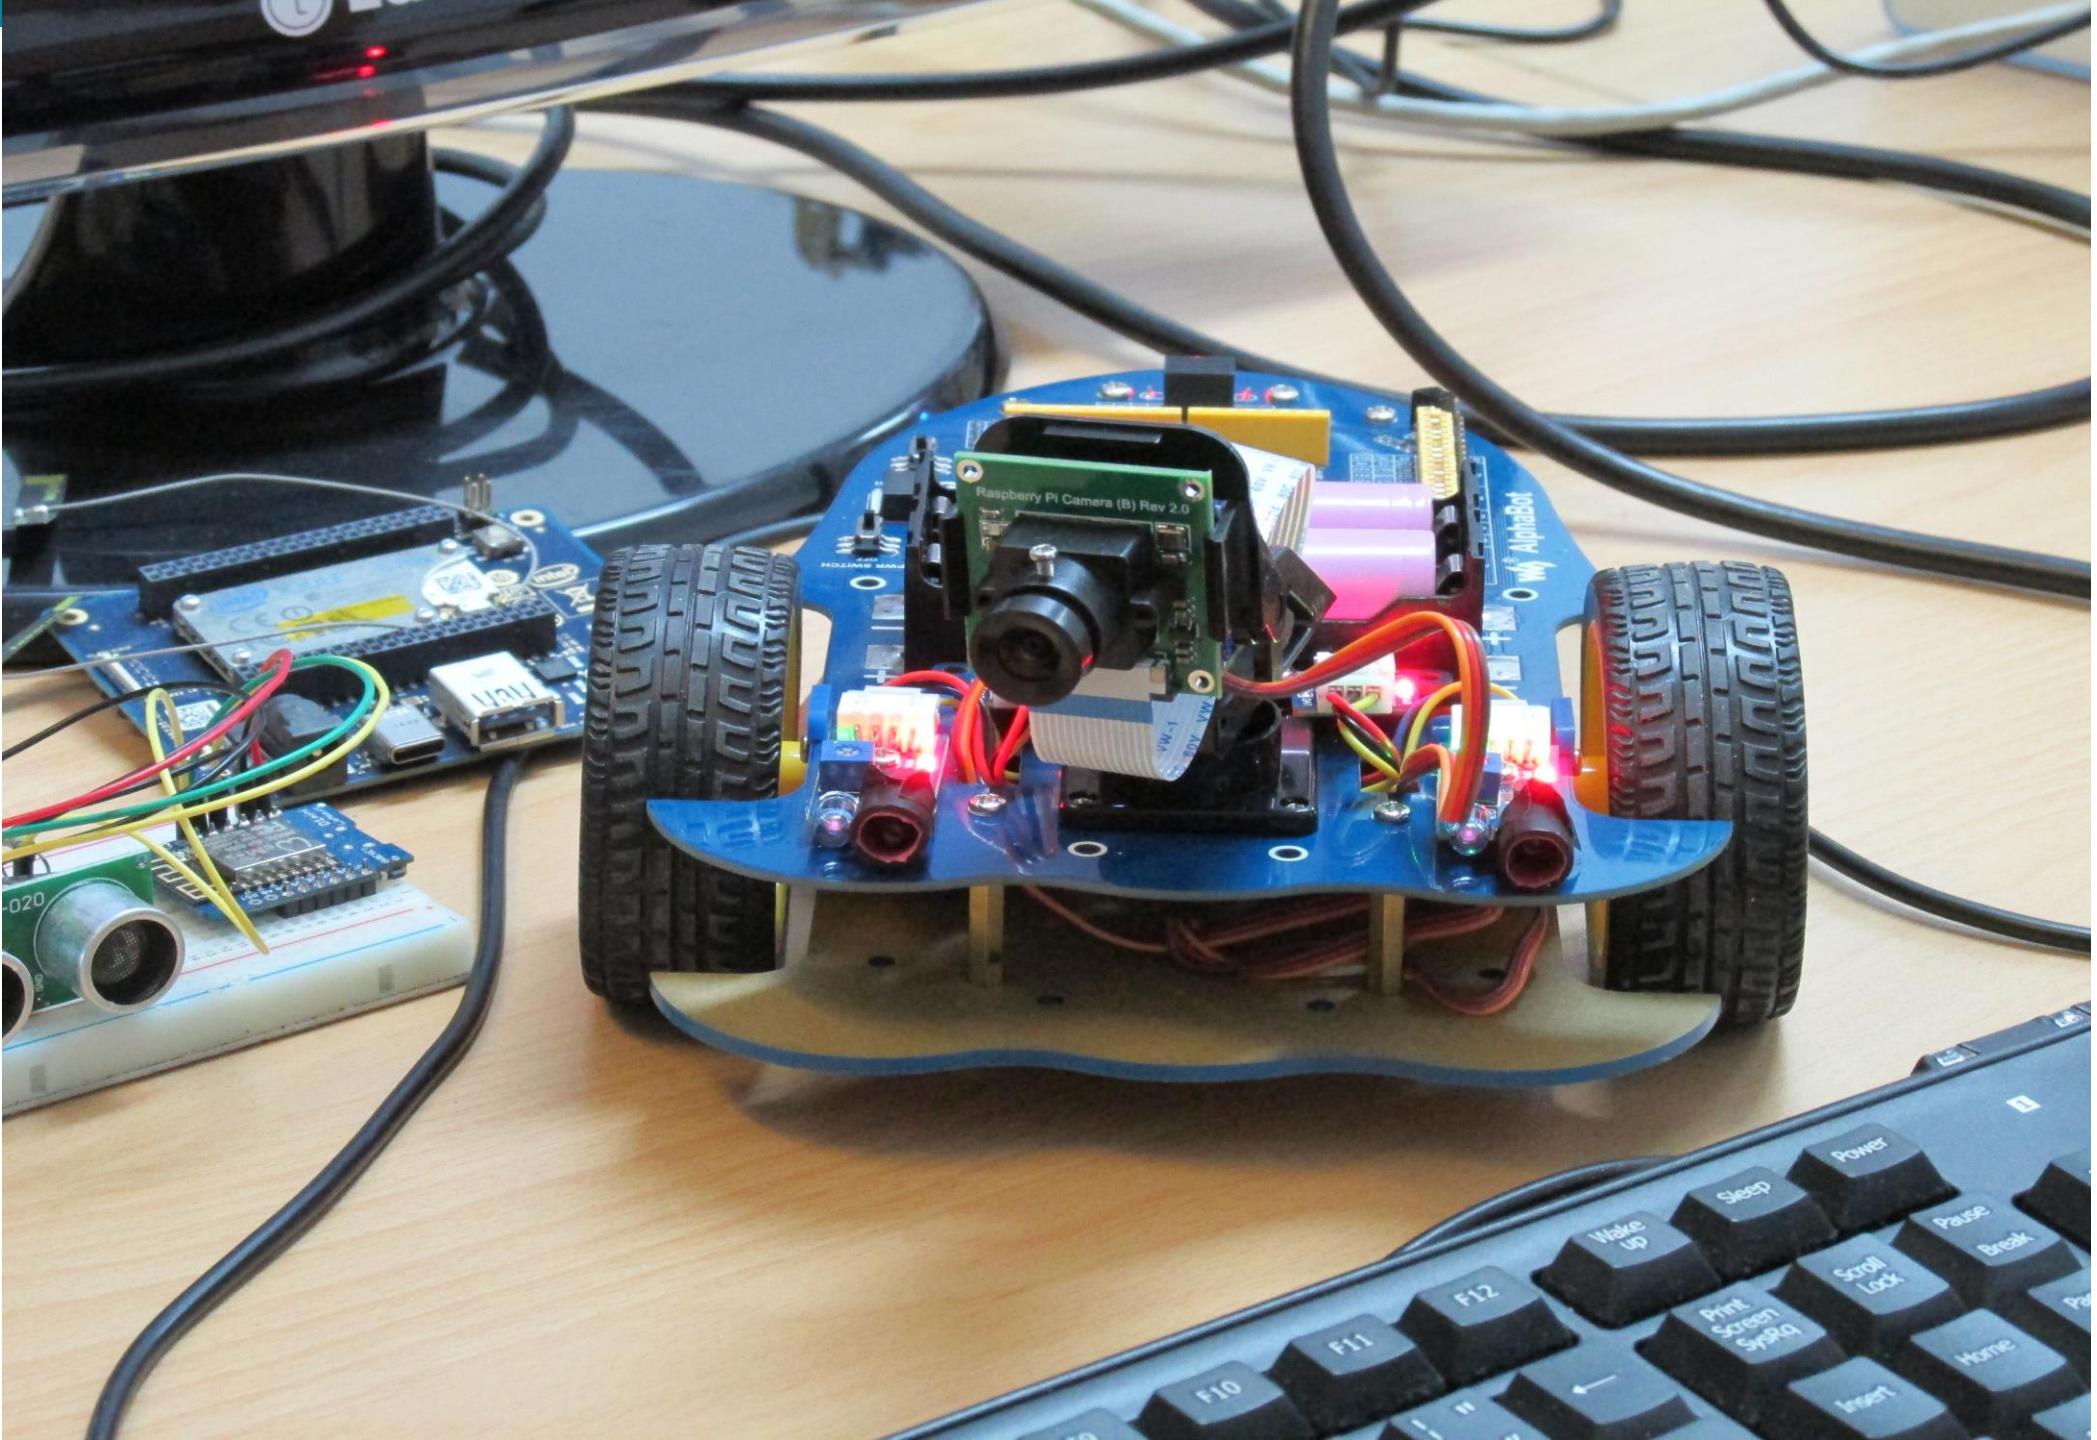


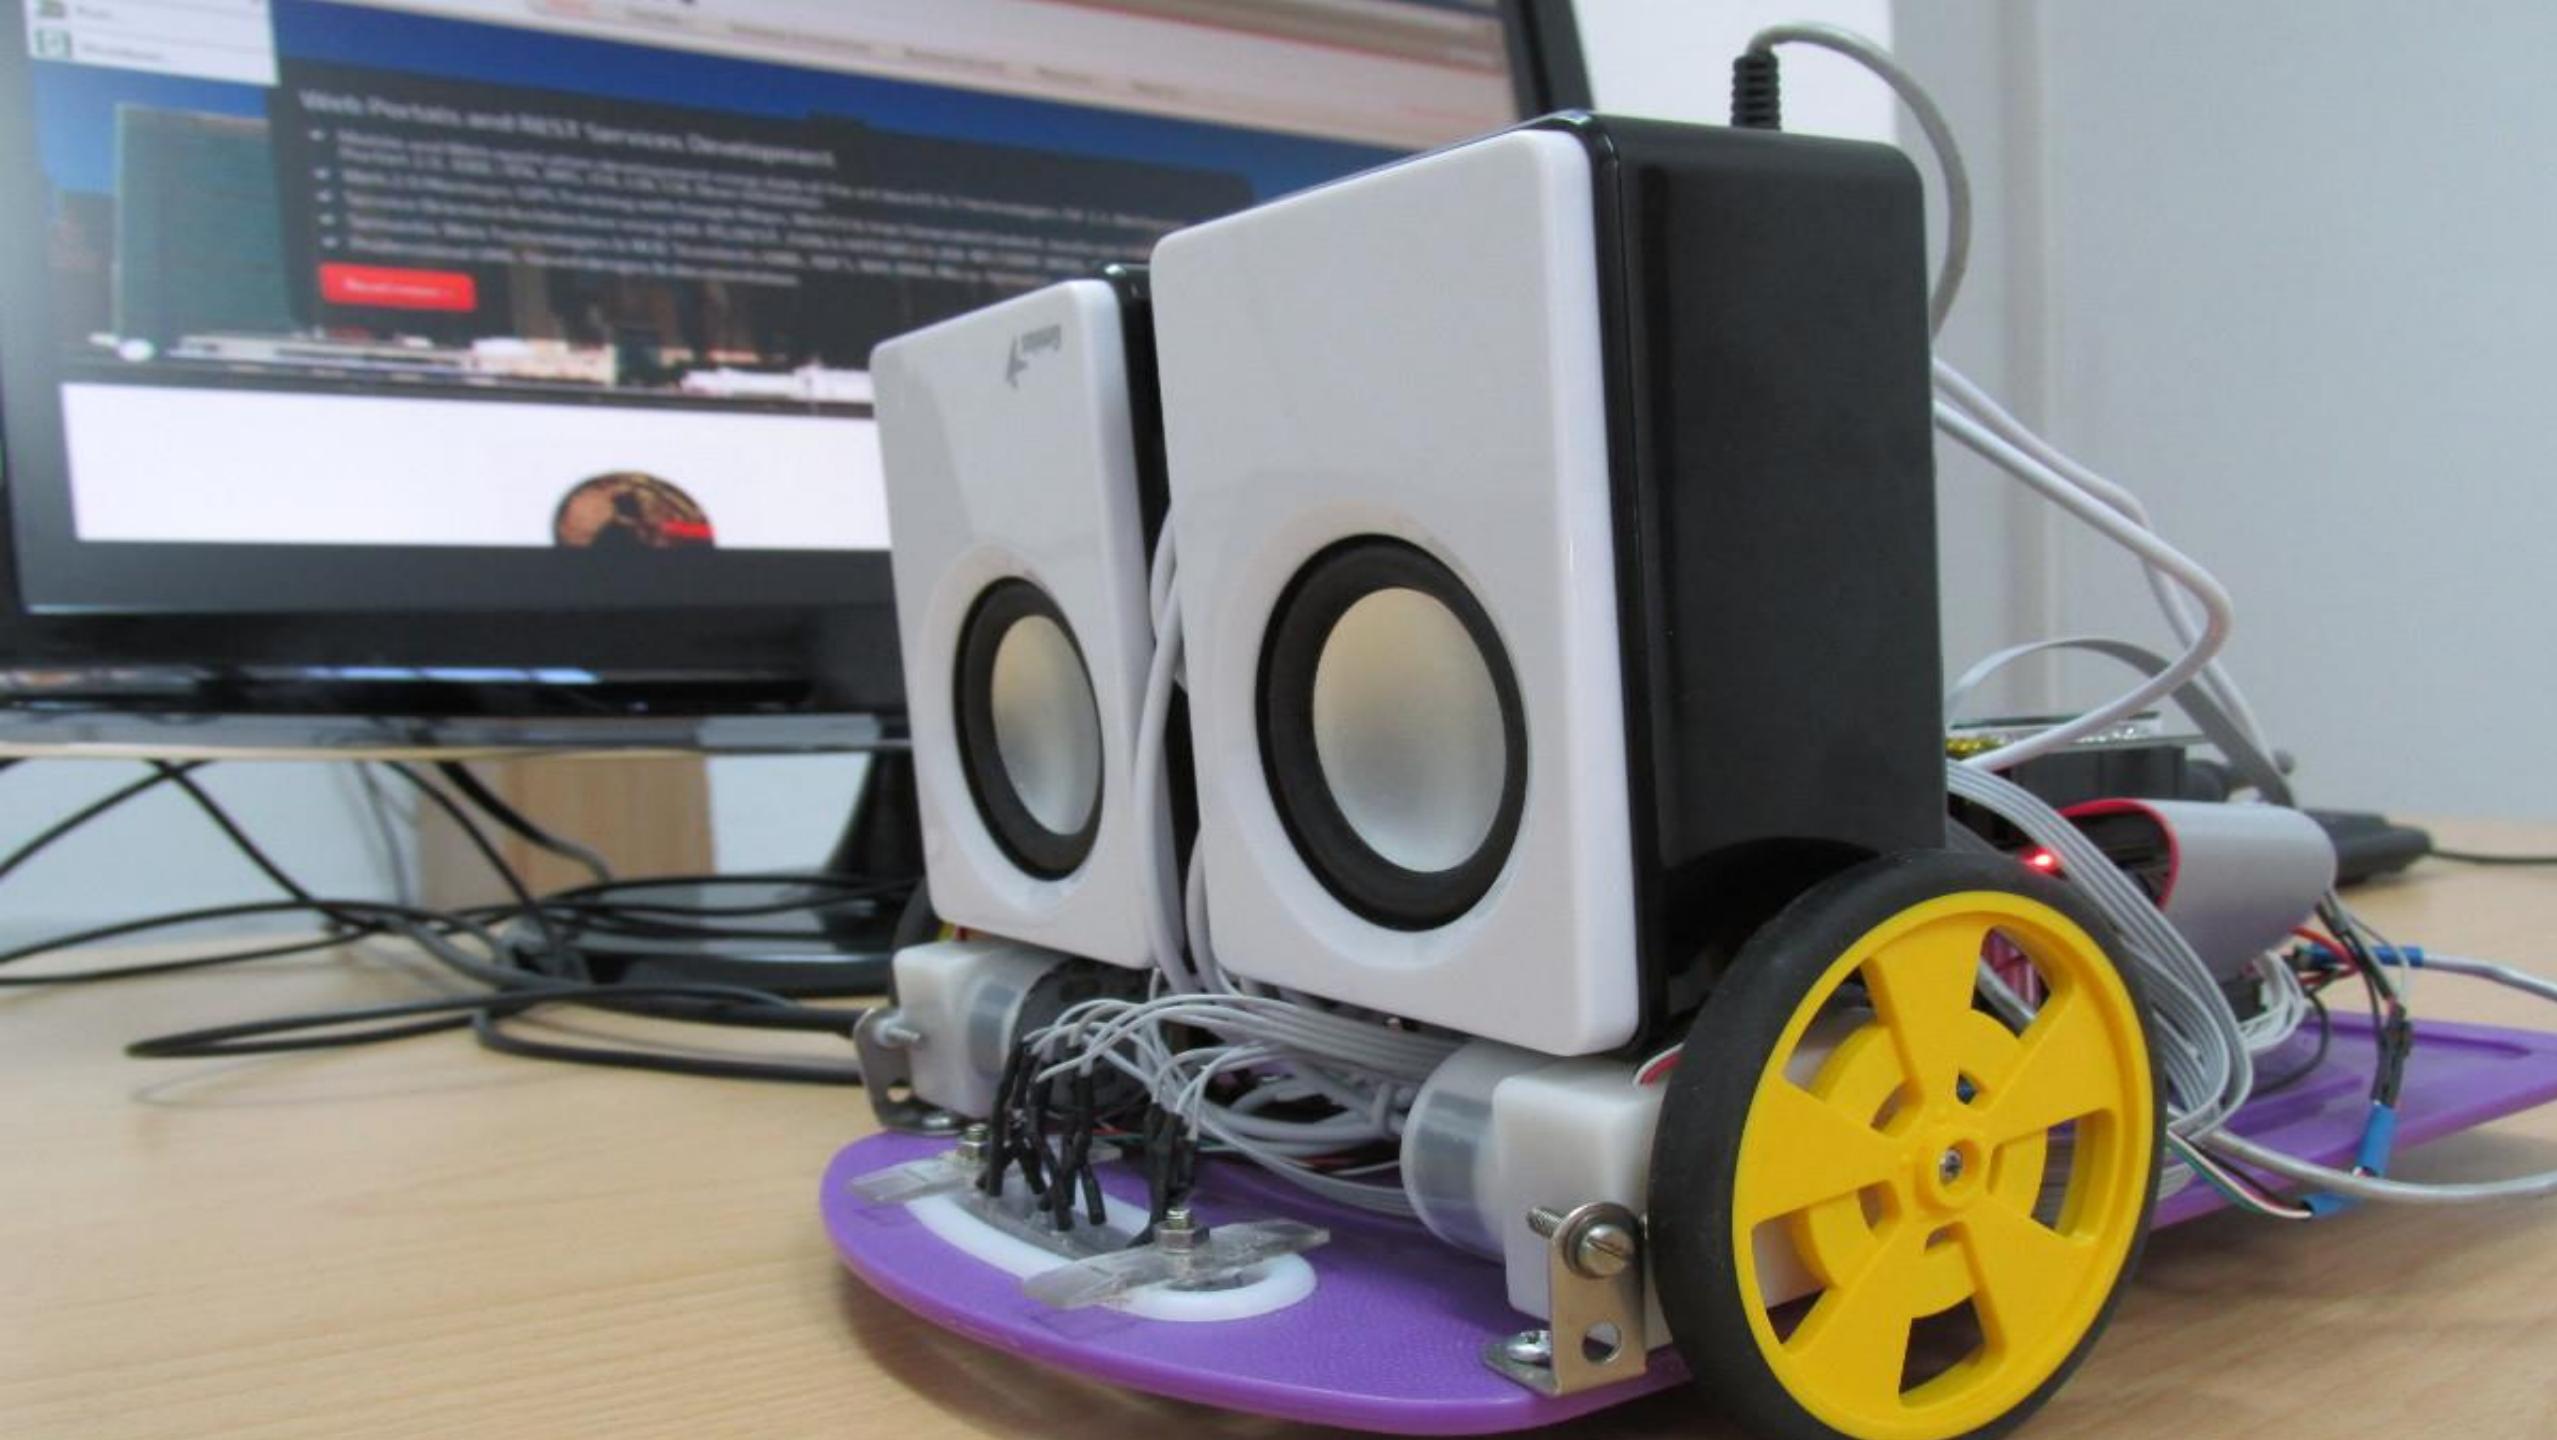
Node Red

Node-RED is a flow-based development tool for visual programming developed originally by IBM for wiring together hardware devices, APIs and online services as part of the Internet of Things.

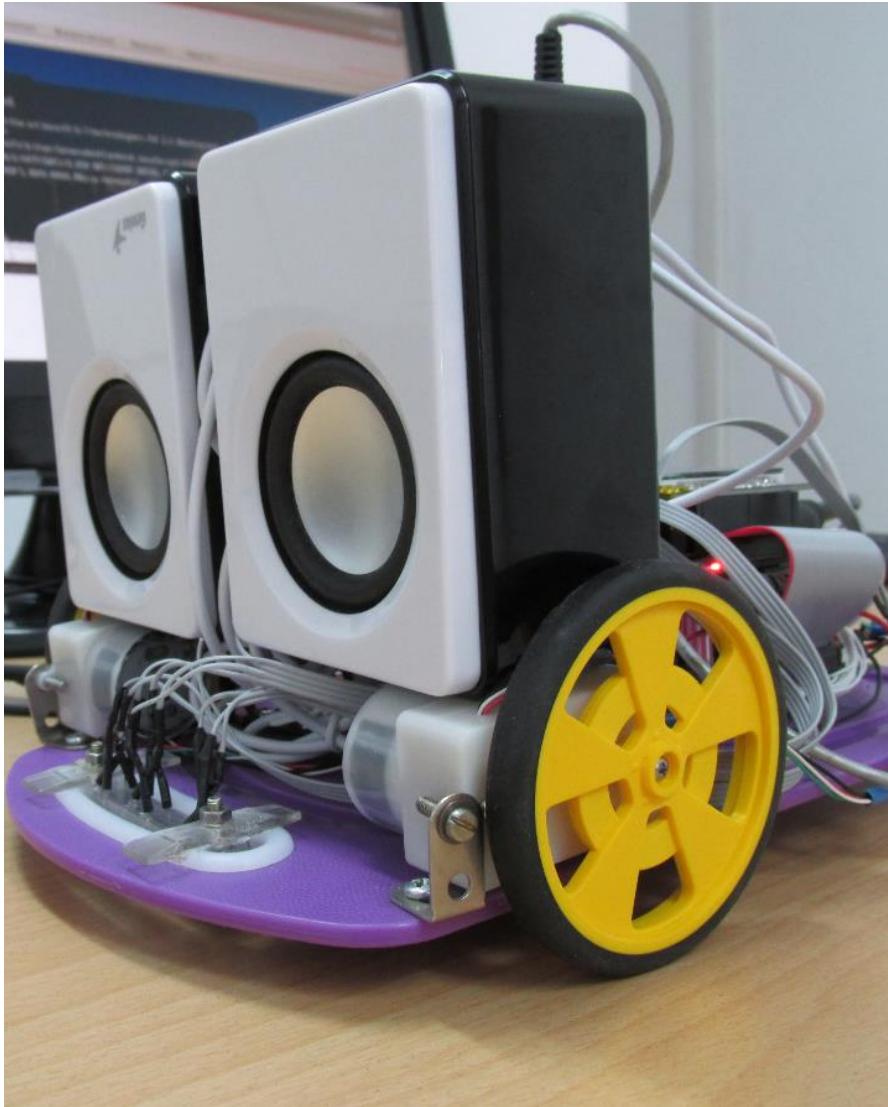








IPTPI: Raspberry Pi + Arduunio Robot



- Raspberry Pi 2 (quad-core ARMv7 @ 900MHz) + Arduino Leonardo clone [A-Star 32U4 Micro](#)
- *Optical encoders* (custom), IR optical array, 3D accelerometers, gyros, and compass [MinIMU-9 v2](#)
- **IPTPI** is programmed in Python, Java and Go using: [Wiring Pi](#), [Numpy](#), [Pandas](#), [Scikit-learn](#), [Pi4J](#), [Reactor](#), [RxJava](#), [GPIO\(Go\)](#)

IPTPI: Raspberry Pi + Arduunio Robot

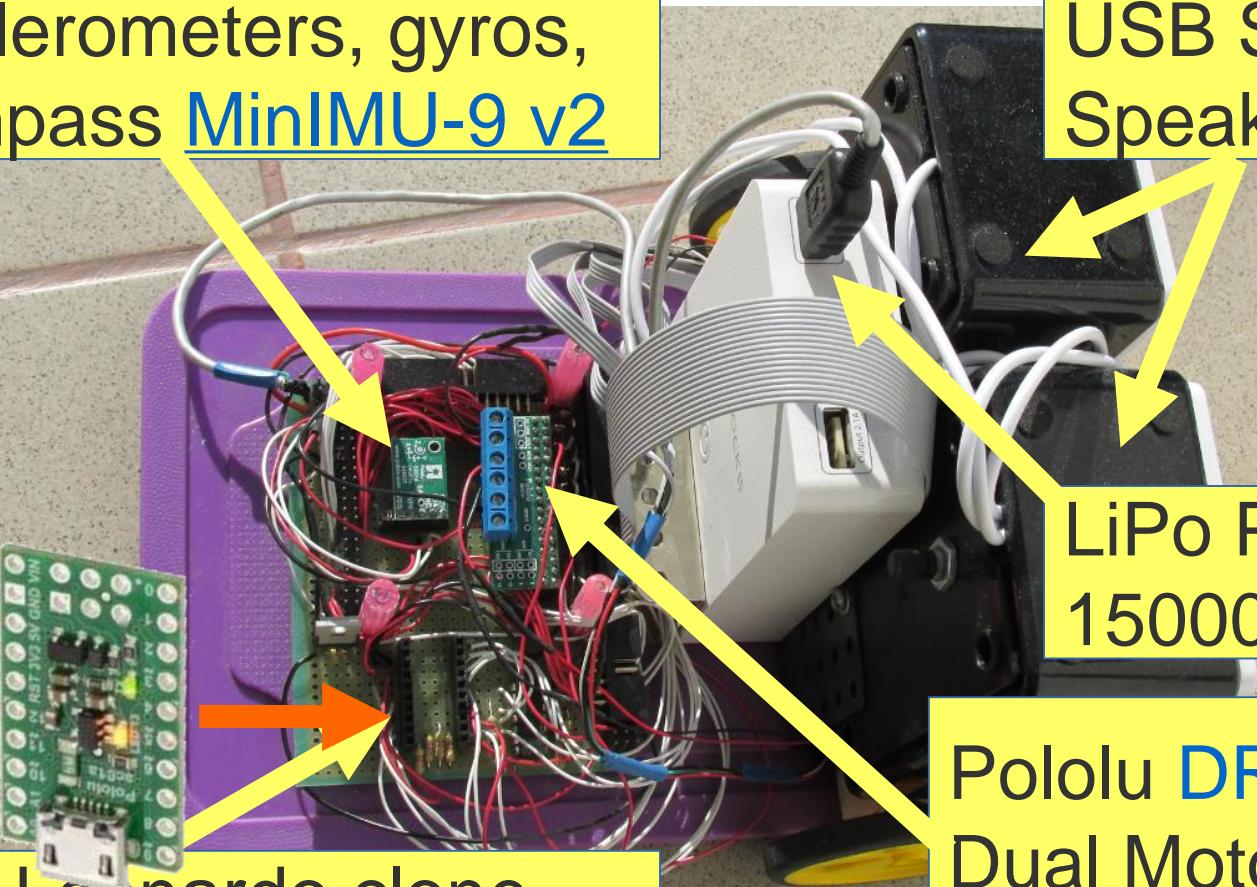
3D accelerometers, gyros,
and compass [MinIMU-9 v2](#)

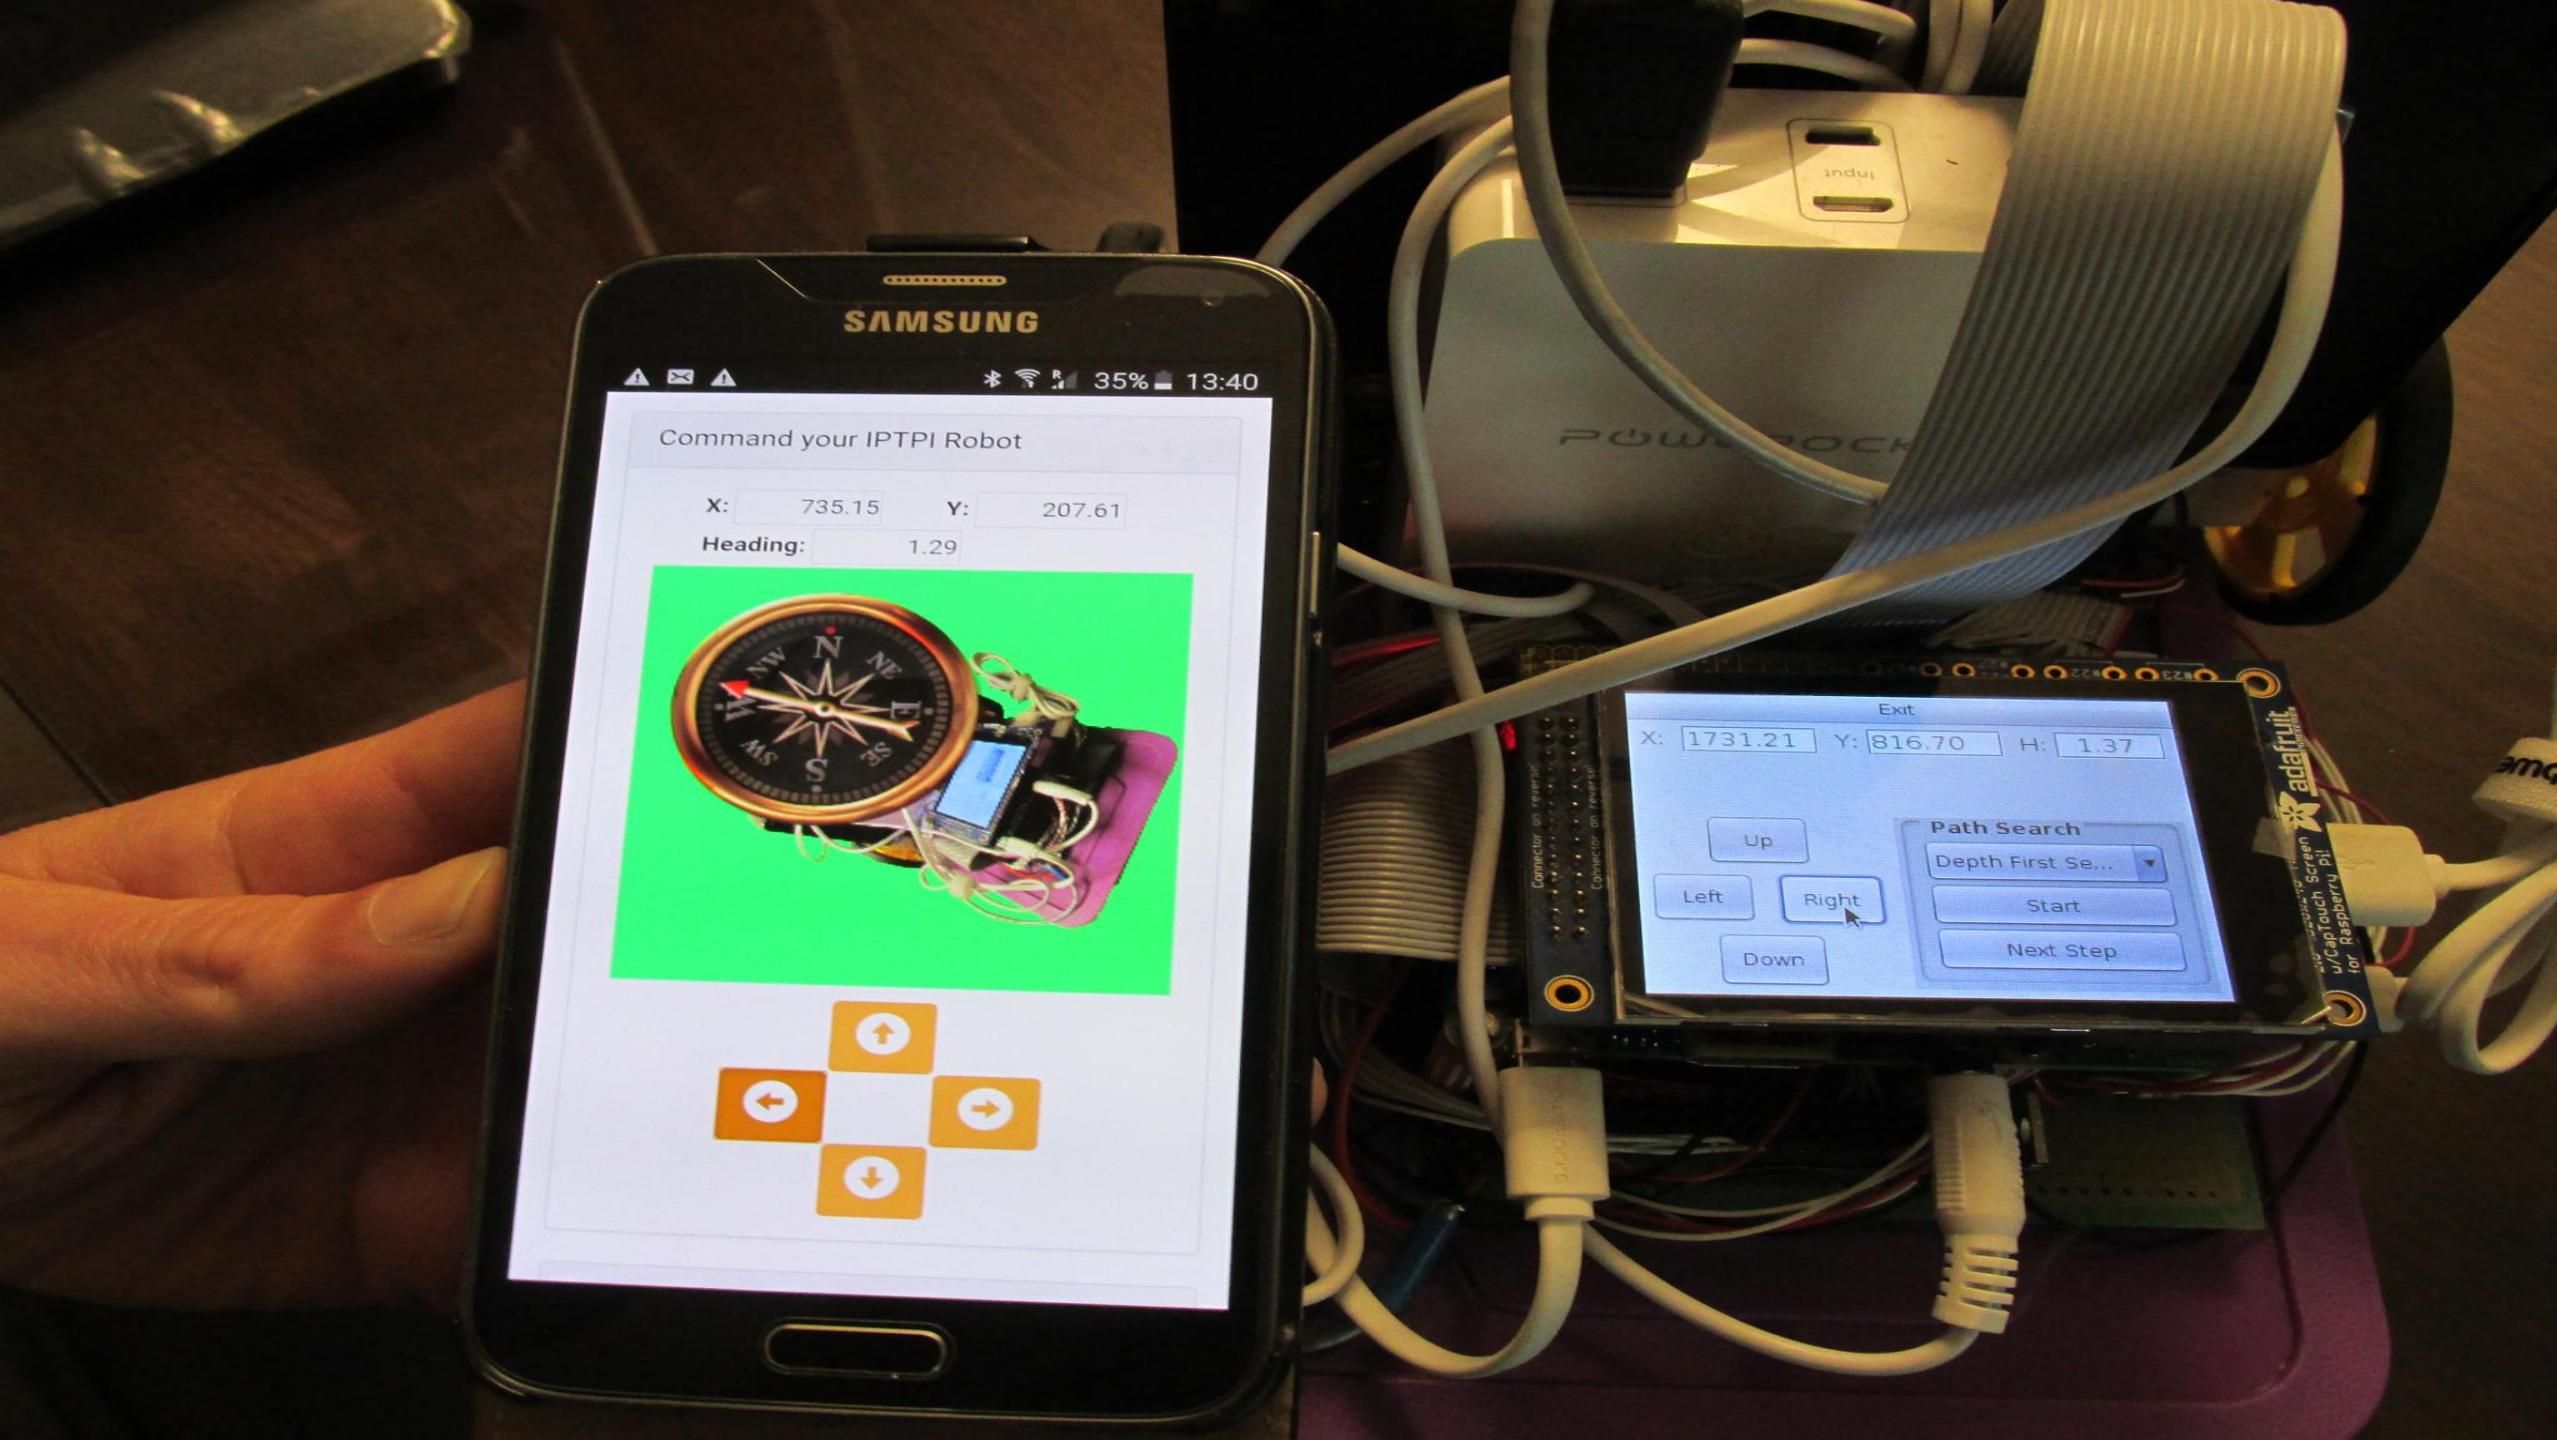
USB Stereo
Speakers - 5V

LiPo Powebank
15000 mAh

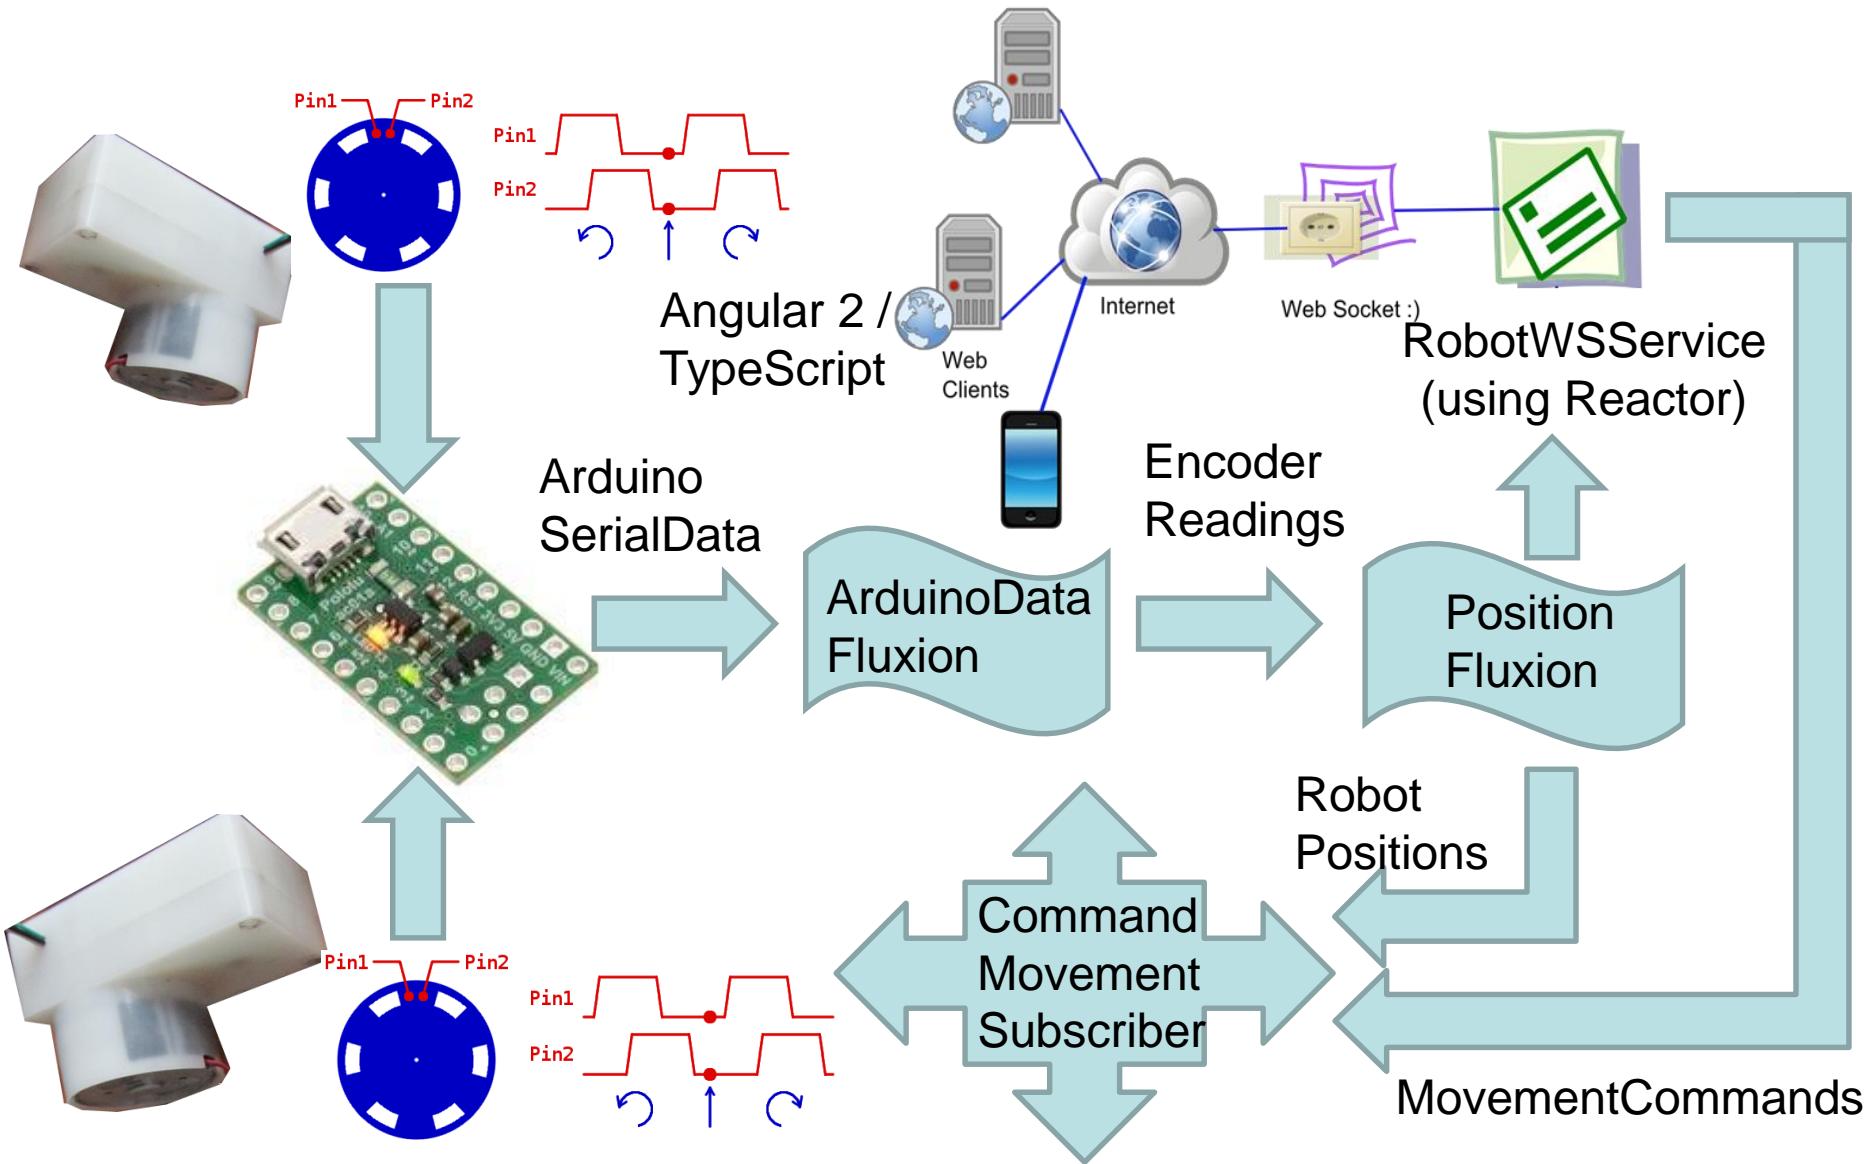
Pololu [DRV8835](#)
Dual Motor Driver
for Raspberry Pi

Arduino Leonardo clone
[A-Star 32U4 Micro](#)





IPTPI Reactive Streams



Functional Reactive Programming

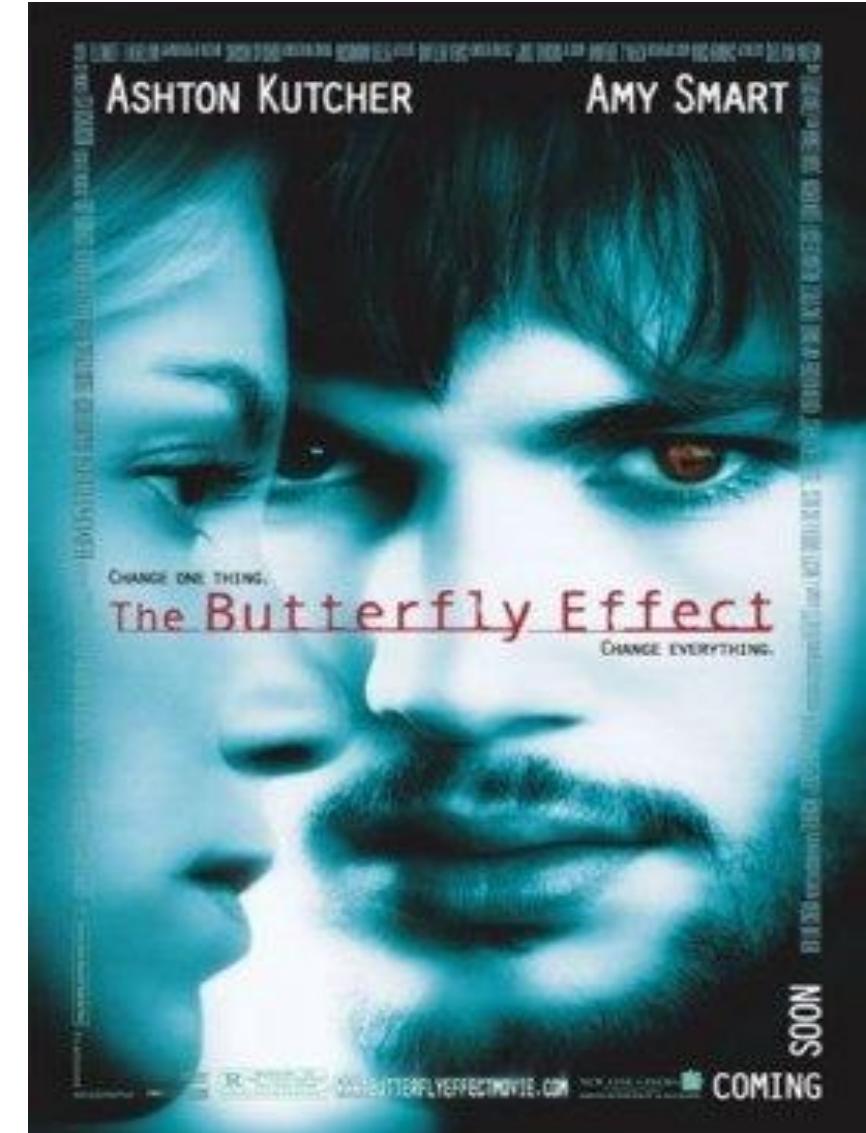


Imperative and Reactive

We live in a Connected Universe

... there is hypothesis that all the things in the Universe are intimately connected, and you can not change a bit without changing all.

Action – Reaction principle is the essence of how Universe behaves.



Imperative vs. Reactive

- **Reactive Programming:** using static or dynamic data flows and propagation of change
- Example: `a := b + c`
- **Functional Programming:** evaluation of mathematical functions, avoids changing-state and mutable data, declarative programming
- Side effects free => much easier to understand and predict the program behavior.

Ex. (RxGo): `observable := rxgo.Just("Hello", "Reactive", "World", "from", "RxGo")().
Map(ToUpper). // map to upper case
Filter(LengthGreaterThan4) // greaterThan4 func filters values > 4
for item := range observable.Observe() {
 fmt.Println(item.V)
}`

Functional Reactive Programming (FRP)

- According to Connal Elliot's (ground-breaking paper at Conference on Functional Programming, 1997), FRP is:

(a) Denotative

(b) Temporally continuous

- FRP is asynchronous data-flow programming using the building blocks of functional programming (map, reduce, filter, etc.) and explicitly modeling time

Reactive Programming

- ReactiveX (Reactive Extensions) - open source polyglot (<http://reactivex.io>):
Rx = Observables + Flow transformations + Schedulers
- Go: RxGo, Kotlin: RxKotlin, Java: RxJava, JavaScript: RxJS, Python: RxPY, C#: Rx.NET, Scala: RxScala, Clojure: RxClojure, C++: RxCpp, Ruby: Rx.rb, Python: RxPY, Groovy: RxGroovy, JRuby: RxJRuby, ...
- Reactive Streams Specification (<http://www.reactive-streams.org/>):
 - Publisher – provider of potentially unbounded number of sequenced elements, according to Subscriber(s) demand (backpressure):
onNext* (onError | onComplete)
 - Subscriber, Subscription
 - Processor = Subscriber + Publisher

ReactiveX: Observable vs. Iterable

Example code showing how similar high-order functions can be applied to an Iterable and an Observable

Iterable

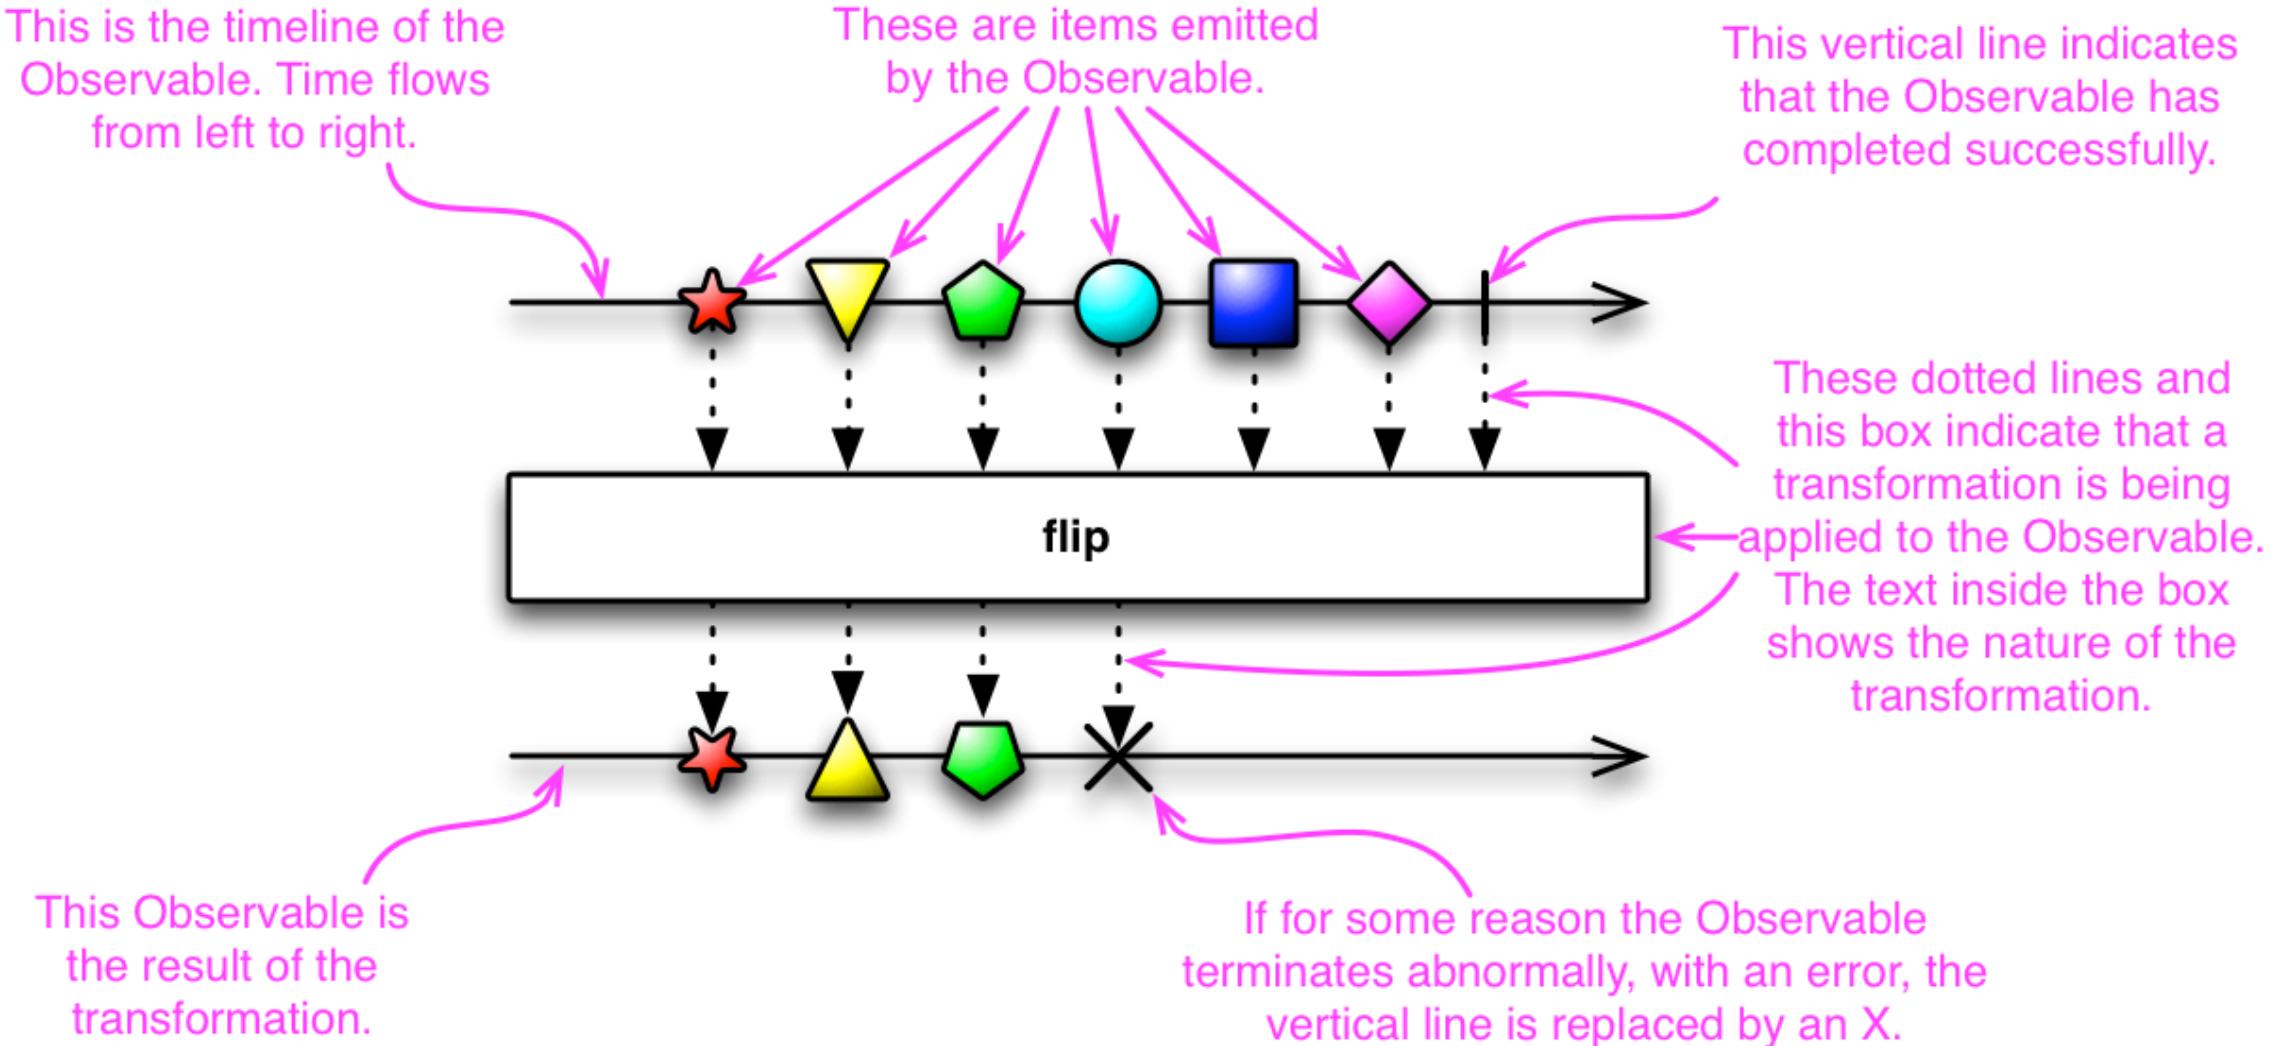
```
getDataFromLocalMemory()  
    .skip(10)  
    .take(5)  
    .map({ s -> return s + " transformed" })  
    .forEach({ println "next => " + it })
```

Observable

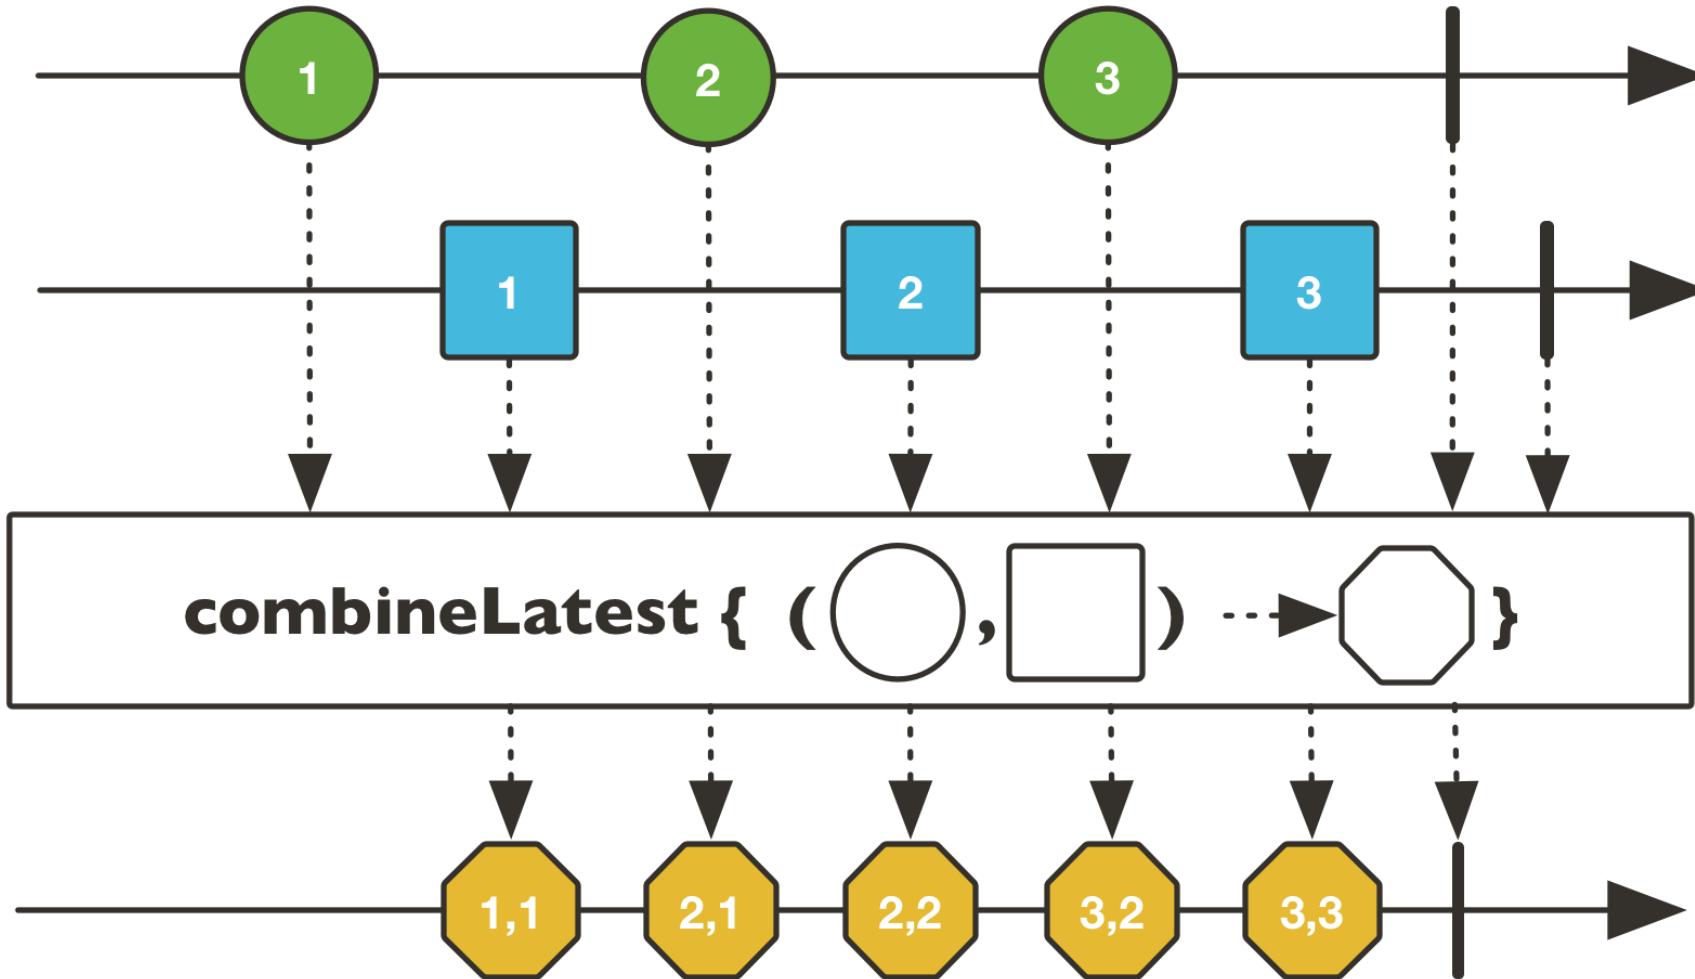
```
getDataFromNetwork()  
    .skip(10)  
    .take(5)  
    .map({ s -> return s + " transformed" })  
    .subscribe({ println "onNext => " + it })
```

You can think of the Observable class as a “**push**” equivalent to [Iterable](#), which is a “**pull**. With an [Iterable](#), the consumer **pulls** values from the producer and the **thread blocks** until those values arrive. By contrast, with an [Observable](#) the producer **pushes** values to the consumer whenever values are available. This approach is more flexible, because **values can arrive synchronously or asynchronously**.

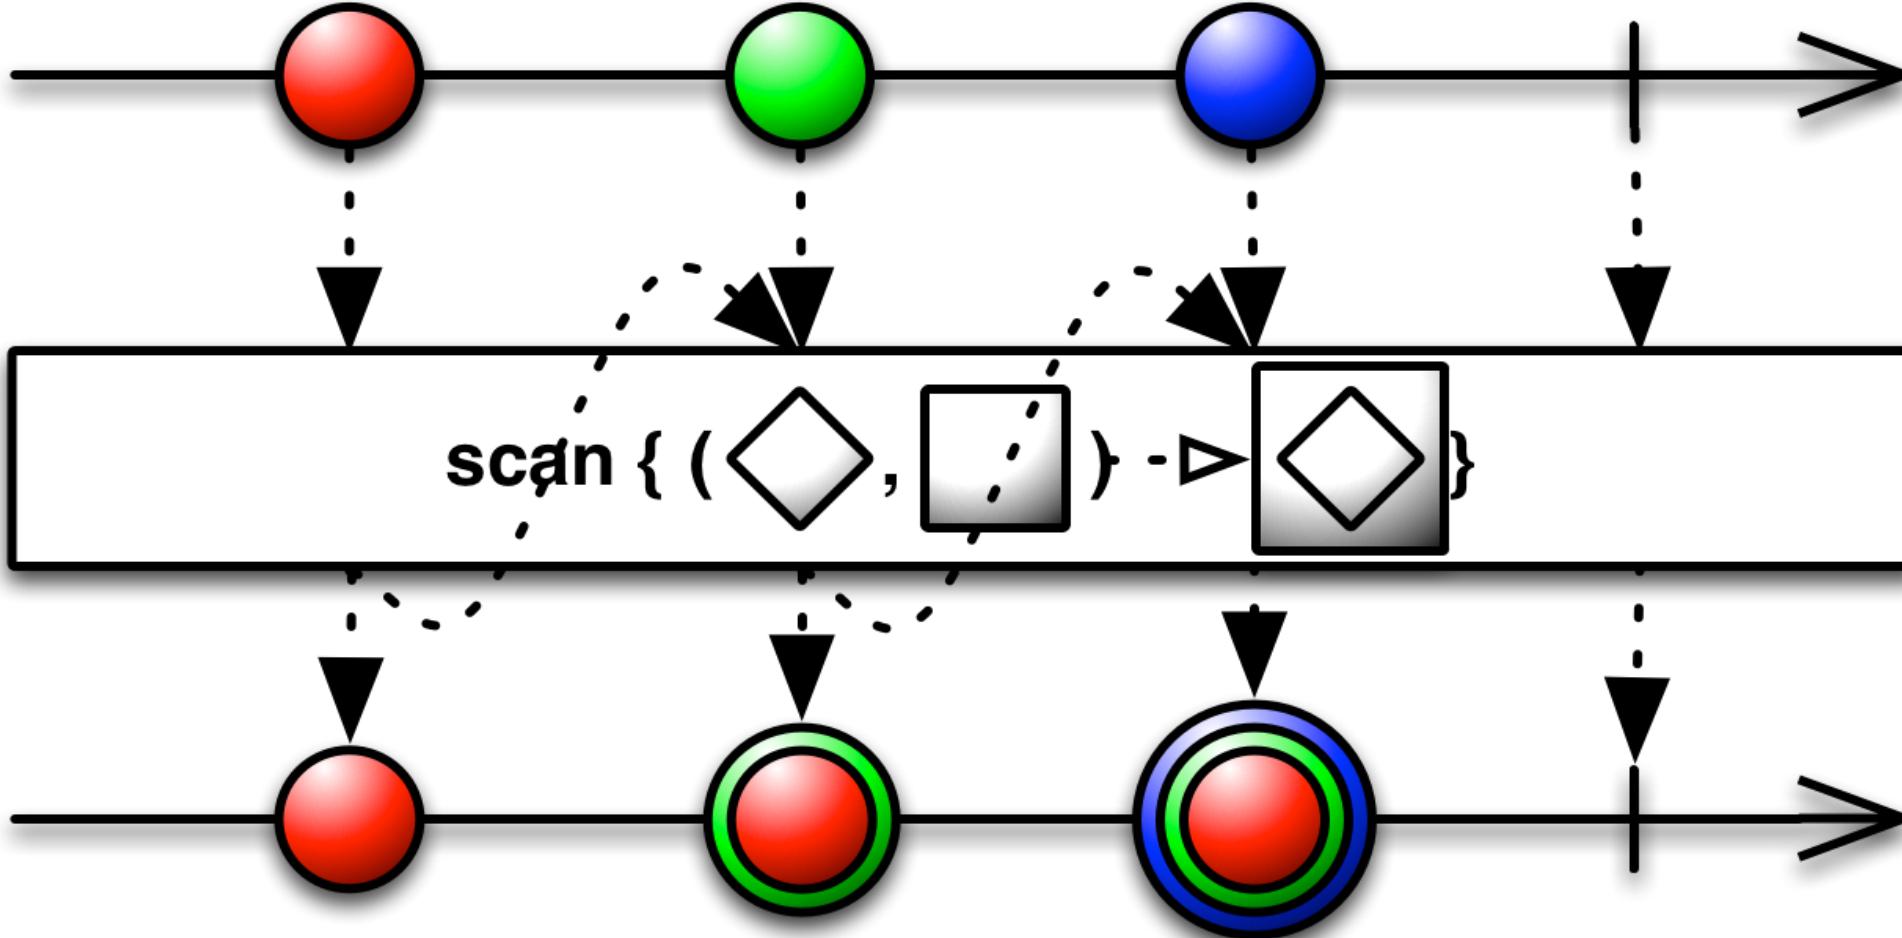
ReactiveX Observable – Marble Diagrams



Example: CombineLatest



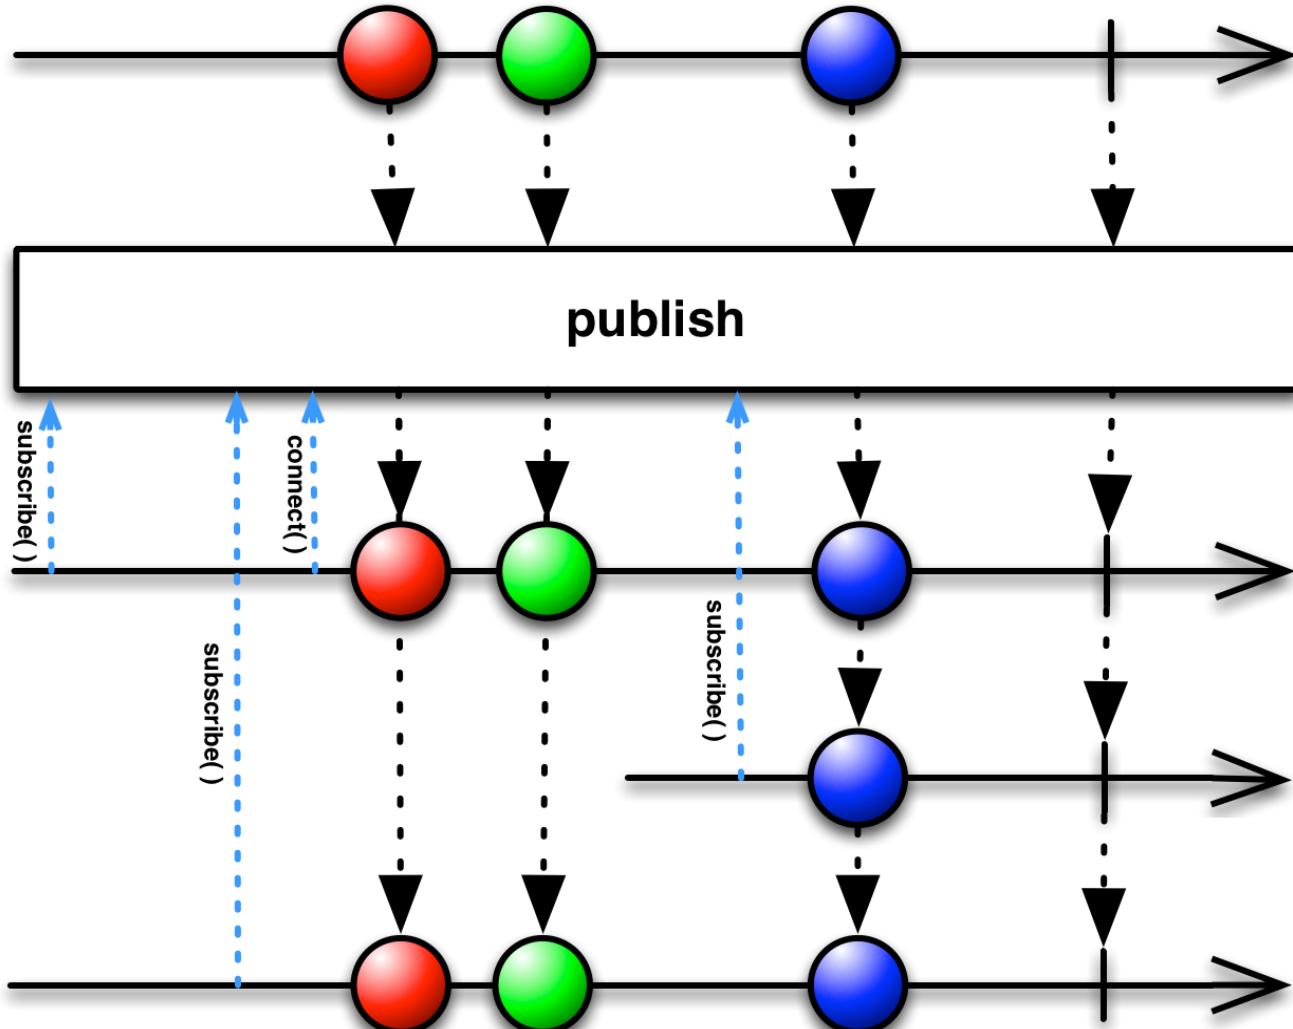
Redux == Rx Scan Operator



Hot and Cold Event Streams

- **PULL-based (Cold Event Streams)** – Cold streams (e.g. RxJava Observable / Flowable or Reactor Flow / Mono) are streams that run their sequence when and if they are subscribed to. They present the sequence from the start to each subscriber.
- **PUSH-based (Hot Event Streams)** – emit values independent of individual subscriptions. They have their own timeline and events occur whether someone is listening or not. When subscription is made observer receives current events as they happen.
- Example: mouse events

Converting Cold to Hot Stream



Reactive Programming

- ❖ Microsoft[®] opens source polyglot project **ReactiveX** (Reactive Extensions) [<http://reactivex.io>]:

Rx = Observables + LINQ + Schedulers :)

Java: RxJava, JavaScript: RxJS, C#: Rx.NET, Scala: RxScala, Clojure: RxClojure, C++: RxCpp, Ruby: Rx.rb, Python: RxPY, Groovy: RxGroovy, JRuby: RxJRuby, Kotlin: RxKotlin ...

- ❖ Reactive Streams Specification
[<http://www.reactive-streams.org/>] used by:
- ❖ (Spring) Project Reactor [<http://projectreactor.io/>]
- ❖ Actor Model – Akka (Java, Scala) [<http://akka.io/>]

Reactive Streams Spec.

- ❖ Reactive Streams – provides standard for **asynchronous stream processing** with non-blocking back pressure.
- ❖ Minimal set of interfaces, methods and protocols for asynchronous data streams
- ❖ April 30, 2015: has been released version 1.0.0 of **Reactive Streams for the JVM** (Java API, Specification, TCK and implementation examples)
- ❖ Java 9: **java.util.concurrent.Flow**

Reactive Streams Spec.

- ❖ **Publisher** – provider of potentially unbounded number of sequenced elements, according to Subscriber(s) demand.

`Publisher.subscribe(Subscriber) => onSubscribe onNext* (onError | onComplete)?`

- ❖ **Subscriber** – calls **Subscription.request(long)** to receive notifications
- ❖ **Subscription** – one-to-one **Subscriber ↔ Publisher**, request data and cancel demand (allow cleanup).
- ❖ **Processor** = **Subscriber + Publisher**

FRP = Async Data Streams

- ❖ FRP is asynchronous data-flow programming using the building blocks of functional programming (e.g. map, reduce, filter) and explicitly modeling time
- ❖ Used for GUIs, robotics, and music. Example (RxJava):

`Observable.from(`

```
    new String[]{"Reactive", "Extensions", "Java"})
    .take(2).map(s -> s + " : on " + new Date())
    .subscribe(s -> System.out.println(s));
```

Result:

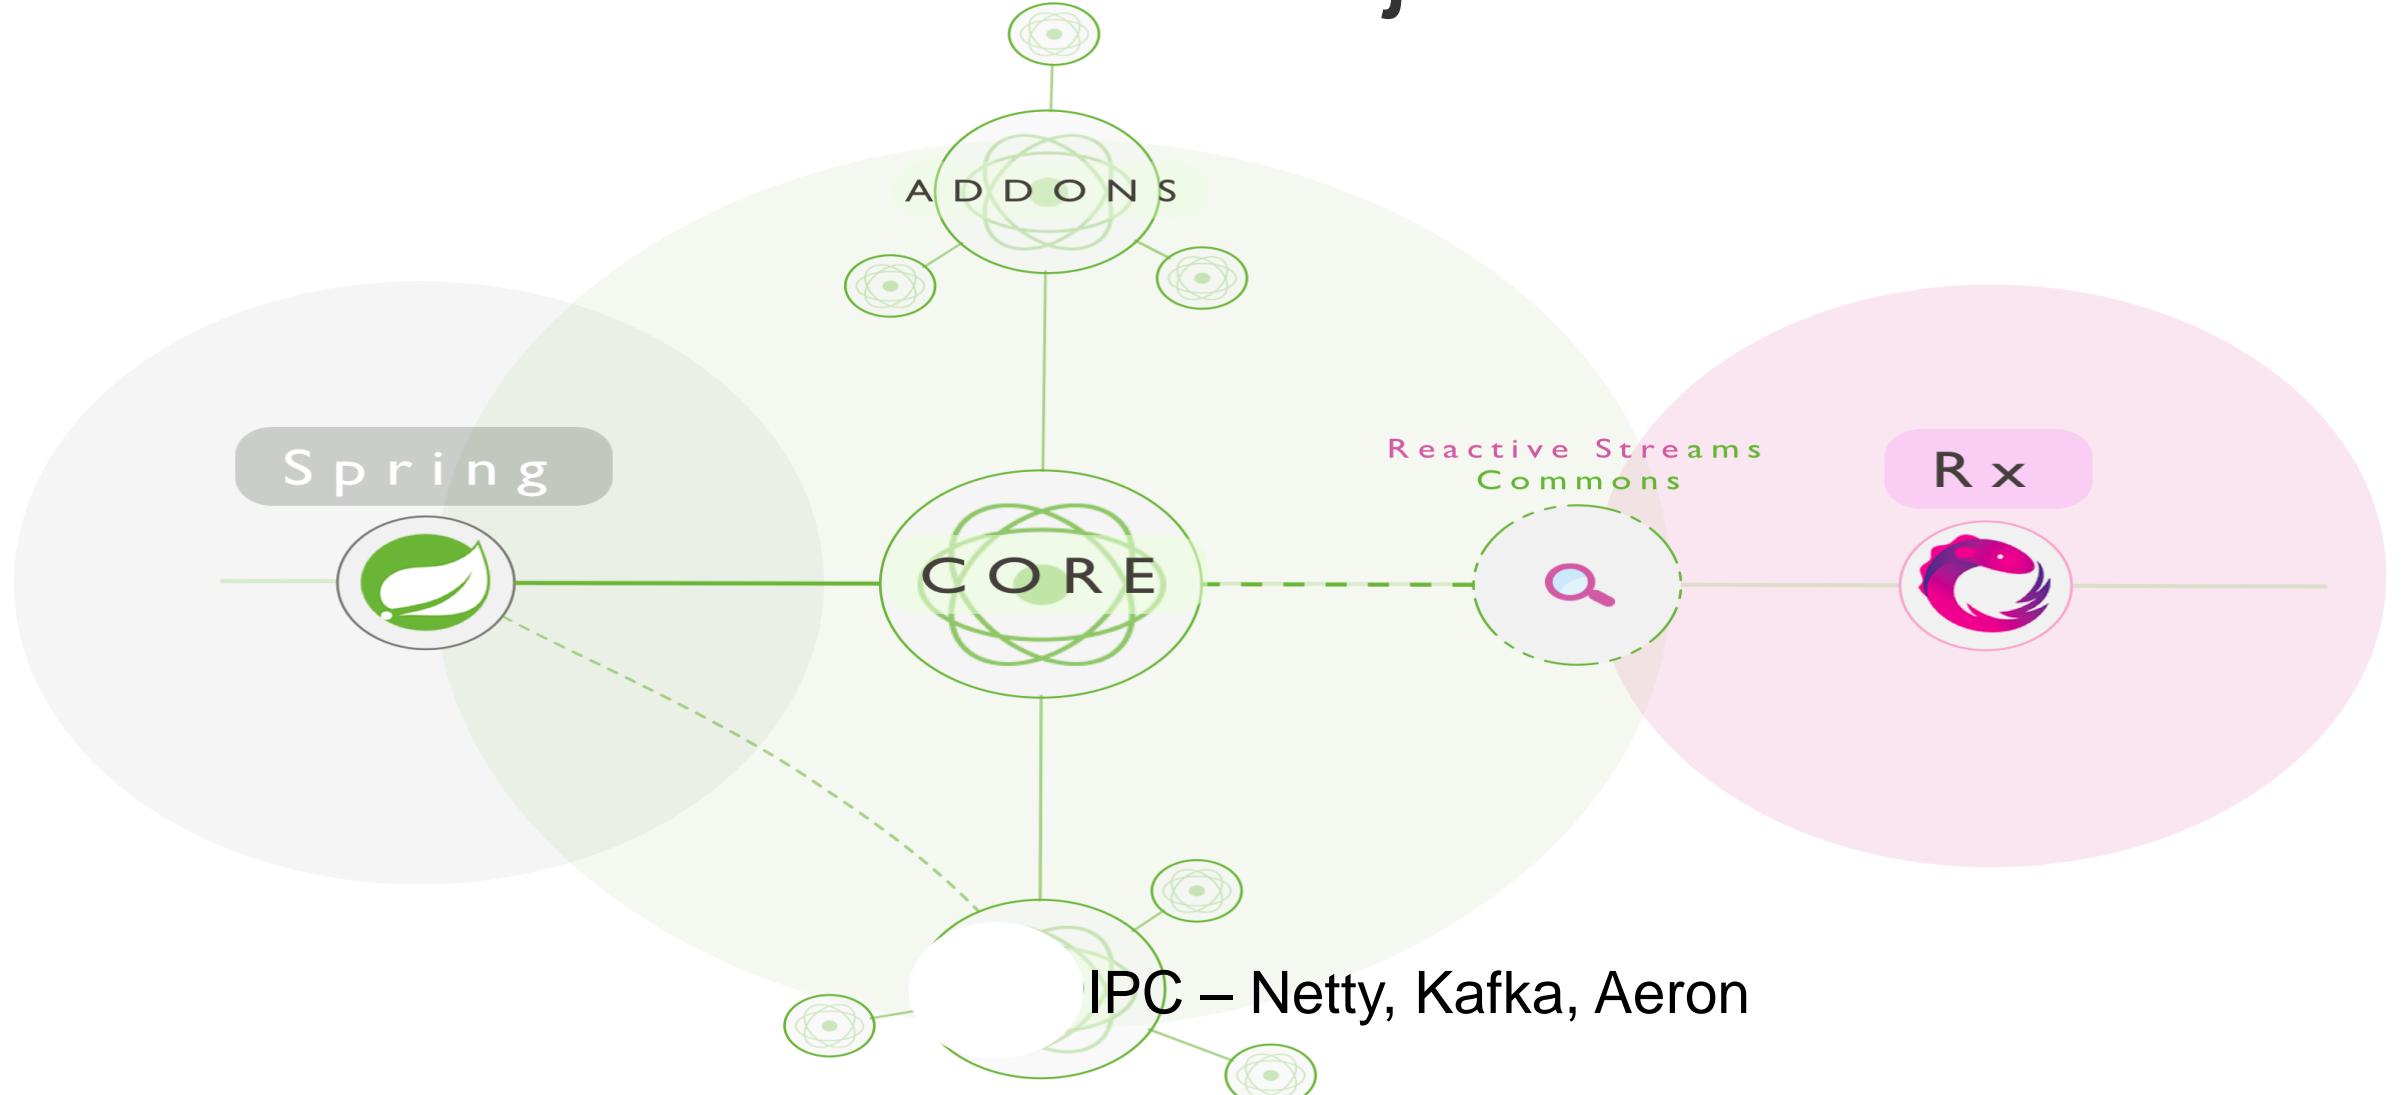
Reactive : on Wed Jun 17 21:54:02 GMT+02:00 2015

Extensions : on Wed Jun 17 21:54:02 GMT+02:00 2015

Project Reactor

- ❖ Reactor project allows building **high-performance (low latency high throughput) non-blocking** asynchronous applications on JVM.
- ❖ Reactor is designed to be extraordinarily fast and can sustain throughput rates on order of **10's of millions of operations per second**.
- ❖ Reactor has powerful API for declaring **data transformations** and **functional composition**.
- ❖ Makes use of the concept of **Mechanical Sympathy** built on top of Disruptor / RingBuffer.

Reactor Projects



<https://github.com/reactor/reactor>, Apache Software License 2.0

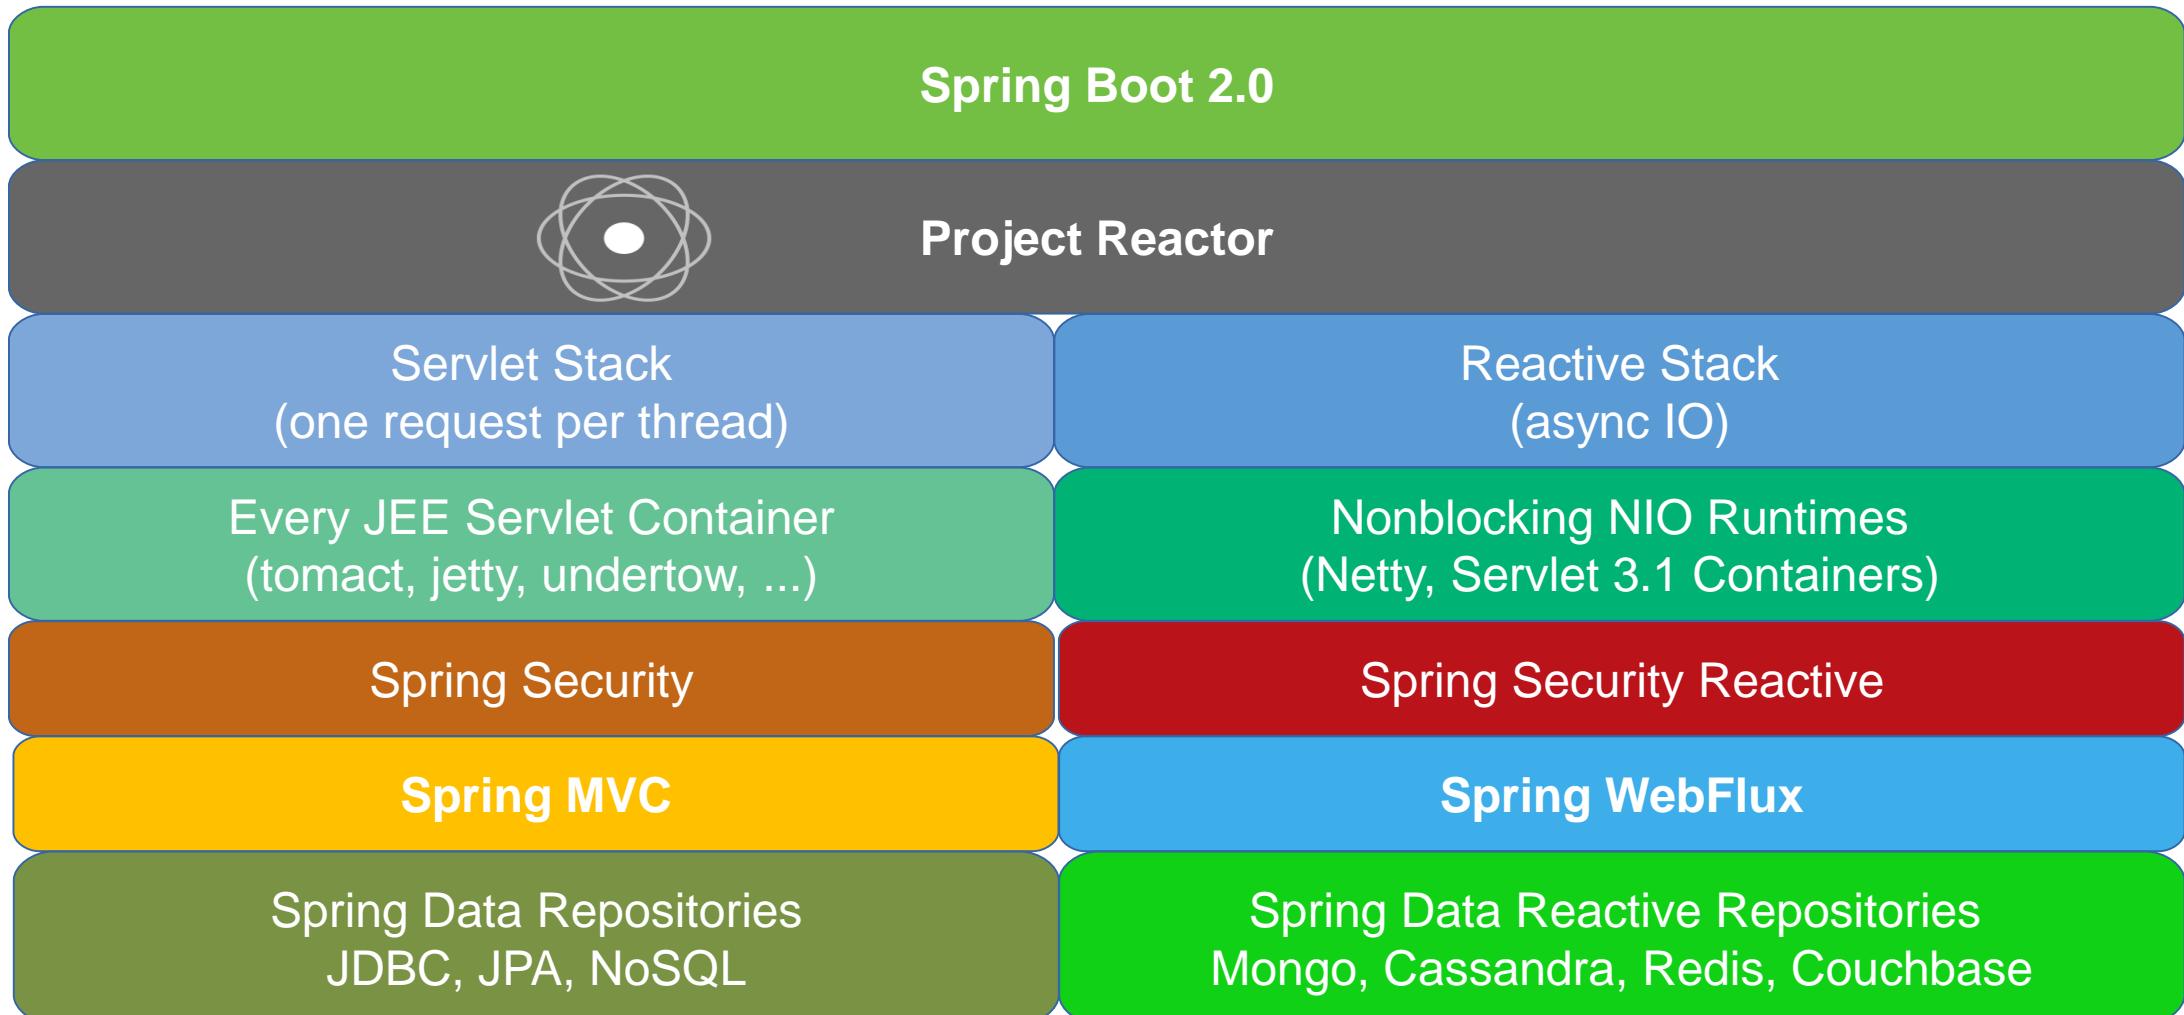
Hot Stream Example - Reactor

```
EmitterProcessor<String> emitter =
    EmitterProcessor.create();
FluxSink<String> sink = emitter.sink();
emitter.publishOn(Schedulers.single())
    .map(String::toUpperCase)
    .filter(s -> s.startsWith("HELLO"))
    .delayElements(Duration.of(1000, MILLIS))
.subscribe(System.out::println);
sink.next("Hello World!"); // emit - non blocking
sink.next("Goodbye World!");
sink.next("Hello Trayan!");
Thread.sleep(3000);
```

Top New Features in Spring 5

- ❖ Reactive Programming Model
- ❖ Spring Web Flux
- ❖ Reactive DB repositories & integrations + hot event streaming:
MongoDB, CouchDB, Redis, Cassandra, Kafka
- ❖ JDK 8+ and Java EE 7+ baseline
- ❖ Testing improvements – WebTestClient (based on reactive WebFlux WebClient)
- ❖ Kotlin functional DSL

Spring 5 Main Building Blocks



Thank's for Your Attention!



Trayan Iliev

IPT – Intellectual Products & Technologies

<http://iproduct.org/>

<https://github.com/iproduct>

<https://twitter.com/trayaniliev>

<https://www.facebook.com/IPT.EACAD>